

Machine Learning techniques to Model Data Intensive Application Performance

A. Battistello P. Ferretti

Contents

| | | |
|----------|---|-----------|
| 1 | Fixed Datasize, all the features | 6 |
| 1.1 | Query R1 | 6 |
| 1.1.1 | Query R1 — Datasize 250GB | 6 |
| 1.1.2 | Query R1 — Datasize 500GB | 7 |
| 1.1.3 | Query R1 — Datasize 750GB | 8 |
| 1.1.4 | Query R1 — Datasize 1000GB | 9 |
| 1.2 | Query R2 | 10 |
| 1.2.1 | Query R2 — Datasize 250GB | 10 |
| 1.2.2 | Query R2 — Datasize 500GB | 11 |
| 1.2.3 | Query R2 — Datasize 750GB | 12 |
| 1.2.4 | Query R2 — Datasize 1000GB | 13 |
| 1.3 | Query R3 | 14 |
| 1.3.1 | Query R3 — Datasize 250GB | 14 |
| 1.3.2 | Query R3 — Datasize 500GB | 15 |
| 1.3.3 | Query R3 — Datasize 750GB | 16 |
| 1.3.4 | Query R3 — Datasize 1000GB | 17 |
| 1.4 | Query R4 | 18 |
| 1.4.1 | Query R4 — Datasize 250GB | 18 |
| 1.4.2 | Query R4 — Datasize 500GB | 19 |
| 1.4.3 | Query R4 — Datasize 750GB | 20 |
| 1.4.4 | Query R4 — Datasize 1000GB | 21 |
| 1.5 | Query R5 | 22 |
| 1.5.1 | Query R5 — Datasize 250GB | 22 |
| 1.5.2 | Query R5 — Datasize 500GB | 23 |
| 1.5.3 | Query R5 — Datasize 750GB | 24 |
| 1.5.4 | Query R5 — Datasize 1000GB | 25 |
| 2 | Fixed Datasize, linear <i>ncores</i>, only <i>ncores</i> | 26 |
| 2.1 | Query R1 | 26 |
| 2.1.1 | Query R1 — Datasize 250GB | 26 |
| 2.1.2 | Query R1 — Datasize 500GB | 27 |
| 2.1.3 | Query R1 — Datasize 750GB | 28 |

| | | |
|----------|---|-----------|
| 2.1.4 | Query R1 — Datasize 1000GB | 29 |
| 2.2 | Query R2 | 30 |
| 2.2.1 | Query R2 — Datasize 250GB | 30 |
| 2.2.2 | Query R2 — Datasize 500GB | 31 |
| 2.2.3 | Query R2 — Datasize 750GB | 32 |
| 2.2.4 | Query R2 — Datasize 1000GB | 33 |
| 2.3 | Query R3 | 34 |
| 2.3.1 | Query R3 — Datasize 250GB | 34 |
| 2.3.2 | Query R3 — Datasize 500GB | 35 |
| 2.3.3 | Query R3 — Datasize 750GB | 36 |
| 2.3.4 | Query R3 — Datasize 1000GB | 37 |
| 2.4 | Query R4 | 38 |
| 2.4.1 | Query R4 — Datasize 250GB | 38 |
| 2.4.2 | Query R4 — Datasize 500GB | 39 |
| 2.4.3 | Query R4 — Datasize 750GB | 40 |
| 2.4.4 | Query R4 — Datasize 1000GB | 41 |
| 2.5 | Query R5 | 42 |
| 2.5.1 | Query R5 — Datasize 250GB | 42 |
| 2.5.2 | Query R5 — Datasize 500GB | 43 |
| 2.5.3 | Query R5 — Datasize 750GB | 44 |
| 2.5.4 | Query R5 — Datasize 1000GB | 45 |
| 2.5.5 | Query R5 — Datasize 1000GB | 45 |
| 3 | Fixed Datasize, $ncores^{-1}$, all the features | 46 |
| 3.1 | Query R1 | 46 |
| 3.1.1 | Query R1 — Datasize 250GB | 46 |
| 3.1.2 | Query R1 — Datasize 500GB | 47 |
| 3.1.3 | Query R1 — Datasize 750GB | 48 |
| 3.1.4 | Query R1 — Datasize 1000GB | 49 |
| 3.2 | Query R2 | 50 |
| 3.2.1 | Query R2 — Datasize 250GB | 50 |
| 3.2.2 | Query R2 — Datasize 500GB | 51 |
| 3.2.3 | Query R2 — Datasize 750GB | 52 |
| 3.2.4 | Query R2 — Datasize 1000GB | 53 |
| 3.3 | Query R3 | 54 |
| 3.3.1 | Query R3 — Datasize 250GB | 54 |
| 3.3.2 | Query R3 — Datasize 500GB | 55 |
| 3.3.3 | Query R3 — Datasize 750GB | 56 |
| 3.3.4 | Query R3 — Datasize 1000GB | 57 |
| 3.4 | Query R4 | 58 |
| 3.4.1 | Query R4 — Datasize 250GB | 58 |
| 3.4.2 | Query R4 — Datasize 500GB | 59 |
| 3.4.3 | Query R4 — Datasize 750GB | 60 |
| 3.4.4 | Query R4 — Datasize 1000GB | 61 |

| | | |
|----------|---|-----------|
| 3.5 | Query R5 | 62 |
| 3.5.1 | Query R5 — Datasize 250GB | 62 |
| 3.5.2 | Query R5 — Datasize 500GB | 63 |
| 3.5.3 | Query R5 — Datasize 750GB | 64 |
| 3.5.4 | Query R5 — Datasize 1000GB | 65 |
| 4 | Fixed Datasize, $ncores^{-1}$, only $ncores$ | 66 |
| 4.1 | Query R1 | 66 |
| 4.1.1 | Query R1 — Datasize 250GB | 66 |
| 4.1.2 | Query R1 — Datasize 500GB | 67 |
| 4.1.3 | Query R1 — Datasize 750GB | 68 |
| 4.1.4 | Query R1 — Datasize 1000GB | 69 |
| 4.2 | Query R2 | 70 |
| 4.2.1 | Query R2 — Datasize 250GB | 70 |
| 4.2.2 | Query R2 — Datasize 500GB | 71 |
| 4.2.3 | Query R2 — Datasize 750GB | 72 |
| 4.2.4 | Query R2 — Datasize 1000GB | 73 |
| 4.3 | Query R3 | 74 |
| 4.3.1 | Query R3 — Datasize 250GB | 74 |
| 4.3.2 | Query R3 — Datasize 500GB | 75 |
| 4.3.3 | Query R3 — Datasize 750GB | 76 |
| 4.3.4 | Query R3 — Datasize 1000GB | 77 |
| 4.4 | Query R4 | 78 |
| 4.4.1 | Query R4 — Datasize 250GB | 78 |
| 4.4.2 | Query R4 — Datasize 500GB | 79 |
| 4.4.3 | Query R4 — Datasize 750GB | 80 |
| 4.4.4 | Query R4 — Datasize 1000GB | 81 |
| 4.5 | Query R5 | 82 |
| 4.5.1 | Query R5 — Datasize 250GB | 82 |
| 4.5.2 | Query R5 — Datasize 500GB | 83 |
| 4.5.3 | Query R5 — Datasize 750GB | 84 |
| 4.5.4 | Query R5 — Datasize 1000GB | 85 |
| 5 | Query comparison | 86 |
| 5.1 | R5 \rightarrow R2 (Fast - Fast) | 86 |
| 5.1.1 | R5 \rightarrow R2 — Datasize 250GB | 86 |
| 5.1.2 | R5 \rightarrow R2 — Datasize 500GB | 87 |
| 5.1.3 | R5 \rightarrow R2 — Datasize 750GB | 88 |
| 5.1.4 | R5 \rightarrow R2 — Datasize 1000GB | 89 |
| 5.2 | R2 \rightarrow R5 (Fast - Fast) | 90 |
| 5.2.1 | R2 \rightarrow R5 — Datasize 250GB | 90 |
| 5.2.2 | R2 \rightarrow R5 — Datasize 500GB | 91 |
| 5.2.3 | R2 \rightarrow R5 — Datasize 750GB | 92 |
| 5.2.4 | R2 \rightarrow R5 — Datasize 1000GB | 93 |

| | | |
|----------|---|------------|
| 5.3 | R3 → R4 (Slow - Slow) | 94 |
| 5.3.1 | R3 → R4 — Datasize 250GB | 94 |
| 5.3.2 | R3 → R4 — Datasize 500GB | 95 |
| 5.3.3 | R3 → R4 — Datasize 750GB | 96 |
| 5.3.4 | R3 → R4 — Datasize 1000GB | 97 |
| 5.4 | R4 → R3 (Slow - Slow) | 98 |
| 5.4.1 | R4 → R3 — Datasize 250GB | 98 |
| 5.4.2 | R4 → R3 — Datasize 500GB | 99 |
| 5.4.3 | R4 → R3 — Datasize 750GB | 100 |
| 5.4.4 | R4 → R3 — Datasize 1000GB | 101 |
| 5.5 | R1, R2, R4 → R3 (Fast and Slow - Average) | 102 |
| 5.5.1 | R1,R2,R4 → R3 — Datasize 250GB | 102 |
| 5.5.2 | R1,R2,R4 → R3 — Datasize 500GB | 103 |
| 5.5.3 | R1,R2,R4 → R3 — Datasize 750GB | 104 |
| 5.5.4 | R1,R2,R4 → R3 — Datasize 1000GB | 105 |
| 5.6 | R2, R5 → R3 (Fast - Slow) | 106 |
| 5.6.1 | R2,R5 → R3 — Datasize 250GB | 106 |
| 5.6.2 | R2,R5 → R3 — Datasize 500GB | 107 |
| 5.6.3 | R2,R5 → R3 — Datasize 750GB | 108 |
| 5.6.4 | R2,R5 → R3 — Datasize 1000GB | 109 |
| 5.7 | R3, R4 → R2 (Fast - Slow) | 110 |
| 5.7.1 | R3,R4 → R2 — Datasize 250GB | 110 |
| 5.7.2 | R3,R4 → R2 — Datasize 500GB | 111 |
| 5.7.3 | R3,R4 → R2 — Datasize 750GB | 112 |
| 5.7.4 | R3,R4 → R2 — Datasize 1000GB | 113 |
| 6 | Fixed Datasize, all the features, | 115 |
| 6.1 | Query Q2 | 115 |
| 6.1.1 | Query Q2 — Datasize 40GB | 115 |
| 6.1.2 | Query Q2 — Datasize 50GB | 116 |
| 6.2 | Query Q3 | 117 |
| 6.2.1 | Query Q3 — Datasize 50GB | 117 |
| 6.3 | Query Q4 | 118 |
| 6.3.1 | Query Q4 — Datasize 50GB | 118 |
| 7 | Fixed Datasize, only nCores, | 119 |
| 7.1 | Query Q2 | 119 |
| 7.1.1 | Query Q2 — Datasize 40GB | 119 |
| 7.1.2 | Query Q2 — Datasize 50GB | 120 |
| 7.2 | Query Q3 | 121 |
| 7.2.1 | Query Q3 — Datasize 50GB | 121 |
| 7.3 | Query Q4 | 122 |
| 7.3.1 | Query Q4 — Datasize 50GB | 122 |

| | | |
|----------|--|------------|
| 8 | Fixed Datasize, all features with 1/nCores, | 123 |
| 8.1 | Query Q2 | 123 |
| 8.1.1 | Query Q2 — Datasize 40GB | 123 |
| 8.1.2 | Query Q2 — Datasize 50GB | 124 |
| 8.2 | Query Q3 | 125 |
| 8.2.1 | Query Q3 — Datasize 50GB | 125 |
| 8.3 | Query Q4 | 126 |
| 8.3.1 | Query Q4 — Datasize 50GB | 126 |
| 9 | Fixed Datasize, only 1/nCores, | 127 |
| 9.1 | Query Q2 | 127 |
| 9.1.1 | Query Q2 — Datasize 40GB | 127 |
| 9.1.2 | Query Q2 — Datasize 50GB | 128 |
| 9.2 | Query Q3 | 129 |
| 9.2.1 | Query Q3 — Datasize 50GB | 129 |
| 9.3 | Query Q4 | 130 |
| 9.3.1 | Query Q4 — Datasize 50GB | 130 |

1 Fixed Datasize, all the features

1.1 Query R1

1.1.1 Query R1 — Datasize 250GB

| Model | RMSE | R^2 | Mean absolute error | Mean relative error |
|--------------------|--------|--------|---------------------|---------------------|
| Linear regression | 0.3041 | 0.9591 | 2480 | 0.0329 |
| Linear SVR | 0.3418 | 0.9830 | 2820 | 0.0375 |
| Polynomial SVR (2) | 0.7566 | 0.7509 | 10643 | 0.2065 |
| Polynomial SVR (3) | 4.2307 | 0.7818 | 23208 | 0.2408 |
| Polynomial SVR (4) | 1.1959 | 0.4622 | 11066 | 0.1907 |
| Polynomial SVR (6) | 5.3772 | 0.4964 | 37535 | 0.4284 |
| Gaussian SVR | 1.2070 | 0.4548 | 7808 | 0.0864 |

Table 1: Results for R1-250




Figure 1: Completion time vs ncores for query R1 with datasize 250GB

1.1.2 Query R1 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0905 | 0.9954 | 5714 | 0.0293 |
| Linear SVR | 0.0994 | 0.9979 | 6740 | 0.0356 |
| Polynomial SVR (2) | 0.4862 | 0.9173 | 30727 | 0.1741 |
| Polynomial SVR (3) | 0.2537 | 0.9662 | 13302 | 0.0643 |
| Polynomial SVR (4) | 0.6772 | 0.8746 | 42865 | 0.2285 |
| Polynomial SVR (6) | 1.3799 | 0.6979 | 67886 | 0.3128 |
| Gaussian SVR | 0.3571 | 0.9779 | 17589 | 0.0756 |

Table 2: Results for R1-500




Figure 2: Completion time vs ncores for query R1 with datasize 500GB

1.1.3 Query R1 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.1813 | 0.9216 | 12361 | 0.0474 |
| Linear SVR | 0.1727 | 0.9390 | 11060 | 0.0413 |
| Polynomial SVR (2) | 0.6893 | 0.1868 | 47136 | 0.1980 |
| Polynomial SVR (3) | 0.4048 | 0.6131 | 27687 | 0.1097 |
| Polynomial SVR (4) | 0.5559 | 0.2879 | 39023 | 0.1624 |
| Polynomial SVR (6) | 0.5809 | 0.2267 | 41711 | 0.1717 |
| Gaussian SVR | 0.1855 | 0.9592 | 12583 | 0.0531 |

Table 3: Results for R1-750




Figure 3: Completion time vs ncores for query R1 with datasize 750GB

1.1.4 Query R1 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.1784 | 0.9416 | 12742 | 0.0332 |
| Linear SVR | 0.1914 | 0.9451 | 13208 | 0.0343 |
| Polynomial SVR (2) | 0.6728 | 0.4426 | 51407 | 0.1411 |
| Polynomial SVR (3) | 0.4131 | 0.7854 | 31112 | 0.0836 |
| Polynomial SVR (4) | 1.0215 | 0.0334 | 55894 | 0.1605 |
| Polynomial SVR (6) | 1.1153 | 0.0093 | 60608 | 0.1735 |
| Gaussian SVR | 0.3595 | 0.8131 | 15070 | 0.0441 |

Table 4: Results for R1-1000




Figure 4: Completion time vs ncores for query R1 with datasize 1000GB

1.2 Query R2

1.2.1 Query R2 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.2914 | 0.9080 | 609 | 0.0073 |
| Linear SVR | 0.2882 | 0.9215 | 622 | 0.0075 |
| Polynomial SVR (2) | 1.0548 | 0.3165 | 2442 | 0.0303 |
| Polynomial SVR (3) | 0.5157 | 0.8003 | 1189 | 0.0149 |
| Polynomial SVR (4) | 0.9650 | 0.1946 | 2218 | 0.0276 |
| Polynomial SVR (6) | 1.0183 | 0.1074 | 2319 | 0.0289 |
| Gaussian SVR | 0.4247 | 0.8933 | 1021 | 0.0126 |

Table 5: Results for R2-250




Figure 5: Completion time vs ncores for query R2 with datasize 250GB

1.2.2 Query R2 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.1427 | 0.9539 | 278 | 0.0039 |
| Linear SVR | 0.1346 | 0.9662 | 279 | 0.0039 |
| Polynomial SVR (2) | 0.6942 | 0.2101 | 1256 | 0.0178 |
| Polynomial SVR (3) | 0.3123 | 0.8255 | 643 | 0.0089 |
| Polynomial SVR (4) | 0.6433 | 0.5642 | 1228 | 0.0174 |
| Polynomial SVR (6) | 0.7042 | 0.4659 | 1277 | 0.0181 |
| Gaussian SVR | 0.2514 | 0.8786 | 574 | 0.0079 |

Table 6: Results for R2-500




Figure 6: Completion time vs ncores for query R2 with datasize 500GB

1.2.3 Query R2 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.2345 | 0.9596 | 469 | 0.0059 |
| Linear SVR | 0.2177 | 0.9748 | 445 | 0.0056 |
| Polynomial SVR (2) | 1.2193 | 0.0000 | 2315 | 0.0292 |
| Polynomial SVR (3) | 0.4706 | 0.8904 | 896 | 0.0114 |
| Polynomial SVR (4) | 1.2193 | 0.0000 | 2315 | 0.0292 |
| Polynomial SVR (6) | 1.2193 | 0.0000 | 2315 | 0.0292 |
| Gaussian SVR | 0.5978 | 0.8044 | 1099 | 0.0138 |

Table 7: Results for R2-750




Figure 7: Completion time vs ncores for query R2 with datasize 750GB

1.2.4 Query R2 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0529 | 0.9971 | 21336 | 0.0227 |
| Linear SVR | 0.0872 | 0.9939 | 39397 | 0.0449 |
| Polynomial SVR (2) | 0.7767 | 0.4473 | 363910 | 0.4851 |
| Polynomial SVR (3) | 0.5370 | 0.8054 | 210031 | 0.2266 |
| Polynomial SVR (4) | 0.9436 | 0.5587 | 387439 | 0.4181 |
| Polynomial SVR (6) | 2.3977 | 0.5408 | 663955 | 0.6566 |
| Gaussian SVR | 0.2309 | 0.9596 | 94043 | 0.1177 |

Table 8: Results for R2-1000




Figure 8: Completion time vs ncores for query R2 with datasize 1000GB

1.3 Query R3

1.3.1 Query R3 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.1109 | 0.9786 | 5118 | 0.0293 |
| Linear SVR | 0.1106 | 0.9889 | 5059 | 0.0292 |
| Polynomial SVR (2) | 0.7172 | 0.1678 | 38067 | 0.2733 |
| Polynomial SVR (3) | 0.2449 | 0.8985 | 13157 | 0.0911 |
| Polynomial SVR (4) | 0.6752 | 0.2430 | 35012 | 0.2465 |
| Polynomial SVR (6) | 0.6198 | 0.3670 | 35378 | 0.2375 |
| Gaussian SVR | 0.1868 | 0.9628 | 8716 | 0.0663 |

Table 9: Results for R3-250




Figure 9: Completion time vs ncores for query R3 with datasize 250GB

1.3.2 Query R3 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0535 | 0.9963 | 11098 | 0.0224 |
| Linear SVR | 0.0606 | 0.9955 | 13859 | 0.0280 |
| Polynomial SVR (2) | 0.2930 | 0.9013 | 59492 | 0.1343 |
| Polynomial SVR (3) | 0.1548 | 0.9691 | 36584 | 0.0771 |
| Polynomial SVR (4) | 0.1646 | 0.9658 | 37111 | 0.0703 |
| Polynomial SVR (6) | 0.3303 | 0.8652 | 60515 | 0.1288 |
| Gaussian SVR | 0.0698 | 0.9937 | 14553 | 0.0288 |

Table 10: Results for R3-500




Figure 10: Completion time vs ncores for query R3 with datasize 500GB

1.3.3 Query R3 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0220 | 0.9996 | 2722 | 0.0034 |
| Linear SVR | 0.0888 | 0.9948 | 13574 | 0.0177 |
| Polynomial SVR (2) | 1.0647 | 0.4290 | 145696 | 0.1758 |
| Polynomial SVR (3) | 0.5284 | 0.8347 | 73215 | 0.0972 |
| Polynomial SVR (4) | 0.8089 | 0.6248 | 113146 | 0.1515 |
| Polynomial SVR (6) | 0.8551 | 0.4785 | 114570 | 0.1545 |
| Gaussian SVR | 0.2530 | 0.9796 | 34529 | 0.0434 |

Table 11: Results for R3-750




Figure 11: Completion time vs ncores for query R3 with datasize 750GB

1.3.4 Query R3 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0734 | 0.9955 | 9574 | 0.0085 |
| Linear SVR | 0.1067 | 0.9922 | 18970 | 0.0187 |
| Polynomial SVR (2) | 0.7430 | 0.6271 | 156805 | 0.1511 |
| Polynomial SVR (3) | 0.3859 | 0.9314 | 63411 | 0.0560 |
| Polynomial SVR (4) | 0.9817 | 0.2514 | 194521 | 0.1816 |
| Polynomial SVR (6) | 1.1743 | 0.0219 | 217951 | 0.2044 |
| Gaussian SVR | 0.3360 | 0.9471 | 54863 | 0.0472 |

Table 12: Results for R3-1000




Figure 12: Completion time vs ncores for query R3 with datasize 1000GB

1.4 Query R4

1.4.1 Query R4 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.1719 | 0.9691 | 5155 | 0.0313 |
| Linear SVR | 0.1665 | 0.9721 | 5476 | 0.0331 |
| Polynomial SVR (2) | 0.8901 | 0.2806 | 37089 | 0.2708 |
| Polynomial SVR (3) | 0.6349 | 0.8391 | 23943 | 0.1357 |
| Polynomial SVR (4) | 0.8922 | 0.4584 | 37940 | 0.2624 |
| Polynomial SVR (6) | 0.8339 | 0.5929 | 35512 | 0.2419 |
| Gaussian SVR | 0.2832 | 0.9483 | 9304 | 0.0558 |

Table 13: Results for R4-250




Figure 13: Completion time vs ncores for query R4 with datasize 250GB

1.4.2 Query R4 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0523 | 0.9941 | 5530 | 0.0157 |
| Linear SVR | 0.0815 | 0.9947 | 13073 | 0.0334 |
| Polynomial SVR (2) | 0.3104 | 0.8825 | 40014 | 0.1194 |
| Polynomial SVR (3) | 0.1446 | 0.9635 | 19321 | 0.0472 |
| Polynomial SVR (4) | 0.2598 | 0.9353 | 35711 | 0.0904 |
| Polynomial SVR (6) | 0.2765 | 0.8948 | 36156 | 0.0950 |
| Gaussian SVR | 0.0816 | 0.9881 | 10981 | 0.0270 |

Table 14: Results for R4-500




Figure 14: Completion time vs ncores for query R4 with datasize 500GB

1.4.3 Query R4 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0215 | 0.9996 | 2345 | 0.0038 |
| Linear SVR | 0.0808 | 0.9957 | 9071 | 0.0162 |
| Polynomial SVR (2) | 0.6955 | 0.7218 | 85402 | 0.1413 |
| Polynomial SVR (3) | 0.3624 | 0.9413 | 39380 | 0.0615 |
| Polynomial SVR (4) | 0.8872 | 0.5977 | 94544 | 0.1441 |
| Polynomial SVR (6) | 1.0199 | 0.3856 | 106583 | 0.1602 |
| Gaussian SVR | 0.1759 | 0.9783 | 14459 | 0.0248 |

Table 15: Results for R4-750




Figure 15: Completion time vs ncores for query R4 with datasize 750GB

1.4.4 Query R4 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0944 | 0.9924 | 56636 | 0.0333 |
| Linear SVR | 0.1024 | 0.9936 | 60674 | 0.0388 |
| Polynomial SVR (2) | 0.8395 | 0.4422 | 580227 | 0.3698 |
| Polynomial SVR (3) | 0.4507 | 0.9060 | 296096 | 0.2340 |
| Polynomial SVR (4) | 0.8560 | 0.6428 | 607746 | 0.4239 |
| Polynomial SVR (6) | 0.9475 | 0.5411 | 673570 | 0.4938 |
| Gaussian SVR | 0.2280 | 0.9589 | 129214 | 0.0695 |

Table 16: Results for R4-1000




Figure 16: Completion time vs ncores for query R4 with datasize 1000GB

1.5 Query R5

1.5.1 Query R5 — Datasize 250GB

| Model | RMSE | R^2 | Mean absolute error | Mean relative error |
|--------------------|--------|---------|---------------------|---------------------|
| Linear regression | 1.1021 | -0.0551 | 713 | 0.0275 |
| Linear SVR | 1.0368 | 0.3190 | 639 | 0.0246 |
| Polynomial SVR (2) | 1.1427 | 0.0277 | 769 | 0.0299 |
| Polynomial SVR (3) | 1.0482 | 0.1294 | 688 | 0.0266 |
| Polynomial SVR (4) | 1.1212 | 0.0921 | 728 | 0.0282 |
| Polynomial SVR (6) | 1.1117 | 0.0337 | 696 | 0.0270 |
| Gaussian SVR | 0.9581 | 0.4034 | 595 | 0.0229 |

Table 17: Results for R5-250




Figure 17: Completion time vs ncores for query R5 with datasize 250GB

1.5.2 Query R5 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.5027 | 0.5247 | 398 | 0.0169 |
| Linear SVR | 0.4869 | 0.6065 | 386 | 0.0163 |
| Polynomial SVR (2) | 0.7132 | 0.2590 | 753 | 0.0318 |
| Polynomial SVR (3) | 0.6277 | 0.4057 | 524 | 0.0221 |
| Polynomial SVR (4) | 0.8572 | 0.0574 | 770 | 0.0325 |
| Polynomial SVR (6) | 0.7629 | 0.0794 | 741 | 0.0314 |
| Gaussian SVR | 0.5268 | 0.5836 | 433 | 0.0182 |

Table 18: Results for R5-500




Figure 18: Completion time vs ncores for query R5 with datasize 500GB

1.5.3 Query R5 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.5755 | 0.6749 | 440 | 0.0176 |
| Linear SVR | 0.6121 | 0.7036 | 465 | 0.0186 |
| Polynomial SVR (2) | 1.4732 | 0.1597 | 1181 | 0.0486 |
| Polynomial SVR (3) | 0.8230 | 0.5766 | 841 | 0.0343 |
| Polynomial SVR (4) | 1.4430 | 0.0262 | 1312 | 0.0534 |
| Polynomial SVR (6) | 1.2012 | 0.0051 | 1189 | 0.0484 |
| Gaussian SVR | 0.6737 | 0.5912 | 521 | 0.0212 |

Table 19: Results for R5-750




Figure 19: Completion time vs ncores for query R5 with datasize 750GB

1.5.4 Query R5 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.5015 | 0.1498 | 1322 | 0.0338 |
| Linear SVR | 0.4946 | 0.3382 | 1208 | 0.0309 |
| Polynomial SVR (2) | 0.4199 | 0.5664 | 947 | 0.0237 |
| Polynomial SVR (3) | 0.5480 | 0.2594 | 1525 | 0.0388 |
| Polynomial SVR (4) | 0.4707 | 0.3918 | 1138 | 0.0287 |
| Polynomial SVR (6) | 0.5688 | 0.3014 | 1456 | 0.0366 |
| Gaussian SVR | 0.3646 | 0.5791 | 820 | 0.0207 |

Table 20: Results for R5-1000




Figure 20: Completion time vs ncores for query R5 with datasize 1000GB

2 Fixed Datasize, linear $ncores$, only $ncores$

2.1 Query R1

2.1.1 Query R1 — Datasize 250GB

| Model | RMSE | R^2 | Mean absolute error | Mean relative error |
|--------------------|--------|--------|---------------------|---------------------|
| Linear regression | 1.2075 | 0.3647 | 11200 | 0.1125 |
| Linear SVR | 1.2849 | 0.4436 | 10934 | 0.1003 |
| Polynomial SVR (2) | 1.5934 | 0.0830 | 22694 | 0.3583 |
| Polynomial SVR (3) | 1.3520 | 0.3512 | 12494 | 0.1275 |
| Polynomial SVR (4) | 1.5463 | 0.0938 | 22056 | 0.3381 |
| Polynomial SVR (6) | 1.5344 | 0.1080 | 22003 | 0.3376 |
| Gaussian SVR | 1.2809 | 0.4200 | 11773 | 0.1226 |

Table 21: Results for R1-250GB, only $ncores$




Figure 21: Completion time vs $ncores$ for query R1 with datasize 250GB

2.1.2 Query R1 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.6040 | 0.6338 | 44332 | 0.3321 |
| Linear SVR | 0.6692 | 0.6428 | 48643 | 0.3623 |
| Polynomial SVR (2) | 0.9736 | 0.0889 | 62658 | 0.5003 |
| Polynomial SVR (3) | 0.5829 | 0.6924 | 40582 | 0.3111 |
| Polynomial SVR (4) | 0.9816 | 0.1055 | 66974 | 0.6097 |
| Polynomial SVR (6) | 0.9863 | 0.1273 | 67430 | 0.6322 |
| Gaussian SVR | 0.4035 | 0.9019 | 30189 | 0.2597 |

Table 22: Results for R1-500GB, only ncores




Figure 22: Completion time vs ncores for query R1 with datasize 500GB

2.1.3 Query R1 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.6470 | 0.7384 | 50197 | 0.2333 |
| Linear SVR | 0.6033 | 0.8114 | 44222 | 0.2349 |
| Polynomial SVR (2) | 1.2098 | 0.1305 | 97577 | 0.4832 |
| Polynomial SVR (3) | 0.6262 | 0.7862 | 46007 | 0.2396 |
| Polynomial SVR (4) | 1.2068 | 0.1404 | 97323 | 0.4831 |
| Polynomial SVR (6) | 1.2002 | 0.1515 | 96594 | 0.4889 |
| Gaussian SVR | 0.5955 | 0.8882 | 43084 | 0.2591 |

Table 23: Results for R1-750GB, only ncores




Figure 23: Completion time vs ncores for query R1 with datasize 750GB

2.1.4 Query R1 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.4444 | 0.8004 | 38407 | 0.0906 |
| Linear SVR | 0.4607 | 0.8045 | 38550 | 0.0889 |
| Polynomial SVR (2) | 0.9613 | 0.1422 | 78118 | 0.1746 |
| Polynomial SVR (3) | 0.4834 | 0.7835 | 36629 | 0.0874 |
| Polynomial SVR (4) | 0.9672 | 0.1294 | 78512 | 0.1755 |
| Polynomial SVR (6) | 0.9746 | 0.1150 | 79059 | 0.1767 |
| Gaussian SVR | 0.2943 | 0.9192 | 18751 | 0.0436 |

Table 24: Results for R1-1000GB, only ncores




Figure 24: Completion time vs ncores for query R1 with datasize 1000GB

2.2 Query R2

2.2.1 Query R2 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.9018 | -0.0459 | 2776 | 0.0338 |
| Linear SVR | 0.9105 | 0.1154 | 2801 | 0.0342 |
| Polynomial SVR (2) | 0.9798 | 0.0107 | 3015 | 0.0368 |
| Polynomial SVR (3) | 0.9001 | 0.1108 | 2702 | 0.0329 |
| Polynomial SVR (4) | 0.9790 | 0.0144 | 3009 | 0.0367 |
| Polynomial SVR (6) | 0.9783 | 0.0199 | 3003 | 0.0366 |
| Gaussian SVR | 0.9095 | 0.1187 | 2800 | 0.0342 |

Table 25: Results for R2-250GB, only ncores




Figure 25: Completion time vs ncores for query R2 with datasize 250GB

2.2.2 Query R2 — Datasize 500GB

| Model | RMSE | R^2 | Mean absolute error | Mean relative error |
|--------------------|--------|---------|---------------------|---------------------|
| Linear regression | 1.2116 | -0.0047 | 2935 | 0.0404 |
| Linear SVR | 1.2260 | 0.0144 | 2980 | 0.0410 |
| Polynomial SVR (2) | 1.2156 | 0.1347 | 2995 | 0.0413 |
| Polynomial SVR (3) | 1.2654 | 0.0000 | 3199 | 0.0443 |
| Polynomial SVR (4) | 1.2183 | 0.1115 | 2985 | 0.0411 |
| Polynomial SVR (6) | 1.2194 | 0.0983 | 2987 | 0.0412 |
| Gaussian SVR | 1.2449 | 0.0064 | 3114 | 0.0431 |

Table 26: Results for R2-500GB, only ncores




Figure 26: Completion time vs ncores for query R2 with datasize 500GB

2.2.3 Query R2 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.9999 | -0.0420 | 2058 | 0.0263 |
| Linear SVR | 1.0055 | 0.0415 | 1982 | 0.0252 |
| Polynomial SVR (2) | 1.0883 | 0.0125 | 2252 | 0.0286 |
| Polynomial SVR (3) | 1.0177 | 0.0581 | 1973 | 0.0251 |
| Polynomial SVR (4) | 1.0924 | 0.0151 | 2266 | 0.0288 |
| Polynomial SVR (6) | 1.1161 | 0.0179 | 2327 | 0.0295 |
| Gaussian SVR | 1.0195 | 0.0150 | 2086 | 0.0266 |

Table 27: Results for R2-750GB, only ncores




Figure 27: Completion time vs ncores for query R2 with datasize 750GB

2.2.4 Query R2 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.5107 | 0.5100 | 255920 | 0.3049 |
| Linear SVR | 0.6277 | 0.6461 | 290435 | 0.3504 |
| Polynomial SVR (2) | 0.7885 | 0.0508 | 318924 | 0.4104 |
| Polynomial SVR (3) | 0.5858 | 0.4917 | 221149 | 0.2926 |
| Polynomial SVR (4) | 0.7500 | 0.0548 | 325104 | 0.3925 |
| Polynomial SVR (6) | 0.7566 | 0.0616 | 329967 | 0.3710 |
| Gaussian SVR | 0.2020 | 0.9521 | 87300 | 0.0985 |

Table 28: Results for R2-1000GB, only ncores




Figure 28: Completion time vs ncores for query R2 with datasize 1000GB

2.3 Query R3

2.3.1 Query R3 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.7631 | 0.4438 | 39680 | 0.1810 |
| Linear SVR | 0.7755 | 0.6918 | 39756 | 0.1747 |
| Polynomial SVR (2) | 1.2251 | 0.0307 | 67244 | 0.3277 |
| Polynomial SVR (3) | 0.7870 | 0.6821 | 40268 | 0.1779 |
| Polynomial SVR (4) | 1.1975 | 0.0394 | 65156 | 0.3207 |
| Polynomial SVR (6) | 1.1661 | 0.0508 | 63994 | 0.3249 |
| Gaussian SVR | 0.7771 | 0.6830 | 41435 | 0.1982 |

Table 29: Results for R3-250GB, only ncores




Figure 29: Completion time vs ncores for query R3 with datasize 250GB

2.3.2 Query R3 — Datasize 500GB

| Model | RMSE | R^2 | Mean absolute error | Mean relative error |
|--------------------|--------|--------|---------------------|---------------------|
| Linear regression | 0.4959 | 0.7518 | 114279 | 0.2329 |
| Linear SVR | 0.5009 | 0.7582 | 114192 | 0.2258 |
| Polynomial SVR (2) | 1.0225 | 0.0258 | 185043 | 0.3254 |
| Polynomial SVR (3) | 0.5107 | 0.7434 | 115022 | 0.2303 |
| Polynomial SVR (4) | 1.0233 | 0.0032 | 189124 | 0.3289 |
| Polynomial SVR (6) | 1.0081 | 0.0024 | 186294 | 0.3216 |
| Gaussian SVR | 0.1438 | 0.9806 | 27885 | 0.0616 |

Table 30: Results for R3-500GB, only ncores




Figure 30: Completion time vs ncores for query R3 with datasize 500GB

2.3.3 Query R3 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.3793 | 0.8368 | 59263 | 0.0822 |
| Linear SVR | 0.3795 | 0.8464 | 58792 | 0.0810 |
| Polynomial SVR (2) | 1.0110 | 0.0321 | 138792 | 0.1985 |
| Polynomial SVR (3) | 0.5158 | 0.7551 | 59479 | 0.0780 |
| Polynomial SVR (4) | 1.0187 | 0.0211 | 139861 | 0.2001 |
| Polynomial SVR (6) | 1.0401 | 0.0000 | 151415 | 0.2135 |
| Gaussian SVR | 0.1659 | 0.9863 | 19628 | 0.0297 |

Table 31: Results for R3-750GB, only ncores




Figure 31: Completion time vs ncores for query R3 with datasize 750GB

2.3.4 Query R3 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.4163 | 0.8159 | 86213 | 0.0920 |
| Linear SVR | 0.4160 | 0.8199 | 86180 | 0.0920 |
| Polynomial SVR (2) | 0.9122 | 0.1475 | 172594 | 0.1799 |
| Polynomial SVR (3) | 0.5750 | 0.6735 | 103081 | 0.1159 |
| Polynomial SVR (4) | 0.9131 | 0.1417 | 172440 | 0.1782 |
| Polynomial SVR (6) | 0.9813 | 0.1347 | 185181 | 0.2003 |
| Gaussian SVR | 0.0971 | 0.9938 | 16264 | 0.0168 |

Table 32: Results for R3-1000GB, only ncores




Figure 32: Completion time vs ncores for query R3 with datasize 1000GB

2.4 Query R4

2.4.1 Query R4 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.7819 | 0.3094 | 32190 | 0.2305 |
| Linear SVR | 0.7724 | 0.3411 | 31610 | 0.2369 |
| Polynomial SVR (2) | 1.1072 | 0.2345 | 43972 | 0.3869 |
| Polynomial SVR (3) | 0.8101 | 0.2724 | 29682 | 0.2136 |
| Polynomial SVR (4) | 1.0995 | 0.1685 | 42974 | 0.3796 |
| Polynomial SVR (6) | 1.0622 | 0.1020 | 42673 | 0.3596 |
| Gaussian SVR | 0.8106 | 0.2932 | 34104 | 0.2553 |

Table 33: Results for R4-250GB, only ncores




Figure 33: Completion time vs ncores for query R4 with datasize 250GB

2.4.2 Query R4 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.5643 | 0.7762 | 91169 | 0.1914 |
| Linear SVR | 0.6193 | 0.8227 | 93111 | 0.1792 |
| Polynomial SVR (2) | 1.3673 | 0.0397 | 168421 | 0.2839 |
| Polynomial SVR (3) | 0.5852 | 0.8088 | 97055 | 0.2142 |
| Polynomial SVR (4) | 1.3574 | 0.0034 | 166877 | 0.2794 |
| Polynomial SVR (6) | 1.3285 | 0.0062 | 162718 | 0.2726 |
| Gaussian SVR | 0.0984 | 0.9942 | 15104 | 0.0406 |

Table 34: Results for R4-500GB, only ncores




Figure 34: Completion time vs ncores for query R4 with datasize 500GB

2.4.3 Query R4 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.4434 | 0.6809 | 53831 | 0.0977 |
| Linear SVR | 0.3957 | 0.7879 | 48222 | 0.0866 |
| Polynomial SVR (2) | 0.7974 | 0.1681 | 84346 | 0.1474 |
| Polynomial SVR (3) | 0.5251 | 0.5563 | 59852 | 0.1091 |
| Polynomial SVR (4) | 0.7986 | 0.1533 | 84417 | 0.1475 |
| Polynomial SVR (6) | 0.8038 | 0.1311 | 84955 | 0.1485 |
| Gaussian SVR | 0.1123 | 0.9797 | 10796 | 0.0194 |

Table 35: Results for R4-750GB, only ncores




Figure 35: Completion time vs ncores for query R4 with datasize 750GB

2.4.4 Query R4 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.5311 | 0.5655 | 342449 | 0.2699 |
| Linear SVR | 0.5613 | 0.5742 | 416843 | 0.3029 |
| Polynomial SVR (2) | 0.9229 | 0.0086 | 713254 | 0.5136 |
| Polynomial SVR (3) | 0.6772 | 0.3007 | 445678 | 0.3277 |
| Polynomial SVR (4) | 0.9151 | 0.0080 | 707727 | 0.5087 |
| Polynomial SVR (6) | 0.9127 | 0.0059 | 705796 | 0.5076 |
| Gaussian SVR | 0.2633 | 0.8936 | 147579 | 0.0877 |

Table 36: Results for R4-1000GB, only ncores




Figure 36: Completion time vs ncores for query R4 with datasize 1000GB

2.5 Query R5

2.5.1 Query R5 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.8259 | -0.3016 | 536 | 0.0214 |
| Linear SVR | 0.7658 | 0.0010 | 512 | 0.0203 |
| Polynomial SVR (2) | 0.7559 | 0.0294 | 503 | 0.0200 |
| Polynomial SVR (3) | 0.7594 | 0.0060 | 504 | 0.0200 |
| Polynomial SVR (4) | 0.7642 | 0.0270 | 507 | 0.0201 |
| Polynomial SVR (6) | 0.7630 | 0.0253 | 508 | 0.0201 |
| Gaussian SVR | 0.8189 | 0.0013 | 486 | 0.0192 |

Table 37: Results for R5-250GB, only ncores




Figure 37: Completion time vs ncores for query R5 with datasize 250GB

2.5.2 Query R5 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 1.1147 | -0.0213 | 1183 | 0.0475 |
| Linear SVR | 1.0929 | 0.1679 | 1192 | 0.0484 |
| Polynomial SVR (2) | 1.1591 | 0.0649 | 1276 | 0.0515 |
| Polynomial SVR (3) | 1.0733 | 0.1551 | 1170 | 0.0476 |
| Polynomial SVR (4) | 1.1591 | 0.0795 | 1276 | 0.0515 |
| Polynomial SVR (6) | 1.1530 | 0.0940 | 1265 | 0.0510 |
| Gaussian SVR | 1.1295 | 0.0016 | 1273 | 0.0518 |

Table 38: Results for R5-500GB, only ncores




Figure 38: Completion time vs ncores for query R5 with datasize 500GB

2.5.3 Query R5 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.8928 | -0.0234 | 1062 | 0.0443 |
| Linear SVR | 1.0199 | -0.0000 | 983 | 0.0425 |
| Polynomial SVR (2) | 1.0199 | -0.0000 | 983 | 0.0425 |
| Polynomial SVR (3) | 1.0199 | -0.0000 | 983 | 0.0425 |
| Polynomial SVR (4) | 1.0199 | -0.0000 | 983 | 0.0425 |
| Polynomial SVR (6) | 1.0199 | -0.0000 | 983 | 0.0425 |
| Gaussian SVR | 1.0199 | -0.0000 | 983 | 0.0425 |

Table 39: Results for R5-750GB, only ncores




Figure 39: Completion time vs ncores for query R5 with datasize 750GB

2.5.4 Query R5 — Datasize 1000GB

2.5.5 Query R5 — Datasize 1000GB

| Model | RMSE | R^2 | Mean absolute error | Mean relative error |
|--------------------|--------|--------|---------------------|---------------------|
| Linear regression | 0.8628 | 0.1111 | 3395 | 0.0850 |
| Linear SVR | 0.8578 | 0.1324 | 3181 | 0.0786 |
| Polynomial SVR (2) | 0.8655 | 0.1903 | 2222 | 0.0509 |
| Polynomial SVR (3) | 0.8833 | 0.1126 | 3533 | 0.0887 |
| Polynomial SVR (4) | 0.8454 | 0.2116 | 2146 | 0.0492 |
| Polynomial SVR (6) | 0.8380 | 0.2376 | 2150 | 0.0493 |
| Gaussian SVR | 0.8023 | 0.3023 | 2506 | 0.0595 |

Table 40: Results for R5-1000GB, only ncores




Figure 40: Completion time vs ncores for query R5 with datasize 1000GB

3 Fixed Datasize, $ncores^{-1}$, all the features

3.1 Query R1

3.1.1 Query R1 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.3569 | 0.9437 | 2586 | 0.0332 |
| Linear SVR | 0.3897 | 0.9727 | 2836 | 0.0362 |
| Polynomial SVR (2) | 0.7267 | 0.7759 | 10787 | 0.2045 |
| Polynomial SVR (3) | 3.6861 | 0.7898 | 20915 | 0.2254 |
| Polynomial SVR (4) | 2.1868 | 0.2291 | 18693 | 0.2484 |
| Polynomial SVR (6) | 3.1348 | 0.3153 | 25360 | 0.3269 |
| Gaussian SVR | 1.2154 | 0.4402 | 7712 | 0.0841 |

Table 41: Results for R1-250GB with the nonlinear $1/ncores$ feature




Figure 41: Completion time vs ncores for query R1 with datasize 250GB

3.1.2 Query R1 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0878 | 0.9957 | 5544 | 0.0283 |
| Linear SVR | 0.0994 | 0.9979 | 6748 | 0.0357 |
| Polynomial SVR (2) | 0.5341 | 0.9077 | 32780 | 0.1738 |
| Polynomial SVR (3) | 0.1464 | 0.9900 | 8855 | 0.0540 |
| Polynomial SVR (4) | 0.6090 | 0.8974 | 38612 | 0.2160 |
| Polynomial SVR (6) | 1.0016 | 0.8146 | 55698 | 0.2804 |
| Gaussian SVR | 0.3567 | 0.9755 | 18202 | 0.0825 |

Table 42: Results for R1-500GB with the nonlinear 1/ncores feature




Figure 42: Completion time vs ncores for query R1 with datasize 500GB

3.1.3 Query R1 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.1409 | 0.9526 | 10604 | 0.0437 |
| Linear SVR | 0.1435 | 0.9626 | 10499 | 0.0427 |
| Polynomial SVR (2) | 0.6469 | 0.2361 | 42774 | 0.1805 |
| Polynomial SVR (3) | 0.4315 | 0.5604 | 27541 | 0.1053 |
| Polynomial SVR (4) | 0.8832 | 0.0817 | 53947 | 0.2200 |
| Polynomial SVR (6) | 0.5945 | 0.2311 | 42316 | 0.1704 |
| Gaussian SVR | 0.1752 | 0.9606 | 11544 | 0.0496 |

Table 43: Results for R1-750GB with the nonlinear 1/ncores feature




Figure 43: Completion time vs ncores for query R1 with datasize 750GB

3.1.4 Query R1 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.1575 | 0.9544 | 10766 | 0.0283 |
| Linear SVR | 0.1449 | 0.9662 | 10234 | 0.0269 |
| Polynomial SVR (2) | 0.6723 | 0.3832 | 49056 | 0.1358 |
| Polynomial SVR (3) | 0.3719 | 0.7768 | 28996 | 0.0774 |
| Polynomial SVR (4) | 0.9733 | 0.0448 | 54529 | 0.1555 |
| Polynomial SVR (6) | 0.9608 | 0.0447 | 55540 | 0.1577 |
| Gaussian SVR | 0.3636 | 0.8120 | 15303 | 0.0446 |

Table 44: Results for R1-1000GB with the nonlinear 1/ncores feature




Figure 44: Completion time vs ncores for query R1 with datasize 1000GB

3.2 Query R2

3.2.1 Query R2 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.3116 | 0.8948 | 656 | 0.0079 |
| Linear SVR | 0.2890 | 0.9187 | 608 | 0.0073 |
| Polynomial SVR (2) | 1.0484 | 0.3489 | 2432 | 0.0302 |
| Polynomial SVR (3) | 0.5203 | 0.7803 | 1223 | 0.0153 |
| Polynomial SVR (4) | 0.9899 | 0.1627 | 2283 | 0.0284 |
| Polynomial SVR (6) | 1.0160 | 0.1256 | 2307 | 0.0287 |
| Gaussian SVR | 0.4203 | 0.8942 | 1005 | 0.0124 |

Table 45: Results for R2-250GB with the nonlinear 1/ncores feature




Figure 45: Completion time vs ncores for query R2 with datasize 250GB

3.2.2 Query R2 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.1459 | 0.9517 | 284 | 0.0040 |
| Linear SVR | 0.1396 | 0.9640 | 293 | 0.0041 |
| Polynomial SVR (2) | 0.6872 | 0.3168 | 1249 | 0.0177 |
| Polynomial SVR (3) | 0.3714 | 0.7477 | 824 | 0.0116 |
| Polynomial SVR (4) | 0.6409 | 0.5494 | 1218 | 0.0172 |
| Polynomial SVR (6) | 0.5920 | 0.5919 | 1350 | 0.0190 |
| Gaussian SVR | 0.2339 | 0.8938 | 536 | 0.0074 |

Table 46: Results for R2-500GB with the nonlinear 1/ncores feature




Figure 46: Completion time vs ncores for query R2 with datasize 500GB

3.2.3 Query R2 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.2407 | 0.9575 | 472 | 0.0060 |
| Linear SVR | 0.2404 | 0.9726 | 476 | 0.0060 |
| Polynomial SVR (2) | 1.2193 | 0.0000 | 2315 | 0.0292 |
| Polynomial SVR (3) | 0.6488 | 0.8600 | 1046 | 0.0133 |
| Polynomial SVR (4) | 1.2193 | 0.0000 | 2315 | 0.0292 |
| Polynomial SVR (6) | 1.2193 | 0.0000 | 2315 | 0.0292 |
| Gaussian SVR | 0.5945 | 0.8089 | 1108 | 0.0139 |

Table 47: Results for R2-750GB with the nonlinear 1/ncores feature




Figure 47: Completion time vs ncores for query R2 with datasize 750GB

3.2.4 Query R2 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0534 | 0.9971 | 21205 | 0.0221 |
| Linear SVR | 0.0843 | 0.9942 | 37707 | 0.0421 |
| Polynomial SVR (2) | 0.7009 | 0.5037 | 336756 | 0.4033 |
| Polynomial SVR (3) | 0.5103 | 0.8111 | 206774 | 0.2177 |
| Polynomial SVR (4) | 0.8812 | 0.5707 | 374286 | 0.4127 |
| Polynomial SVR (6) | 0.8290 | 0.5363 | 351110 | 0.4688 |
| Gaussian SVR | 0.2298 | 0.9590 | 93580 | 0.1148 |

Table 48: Results for R2-1000GB with the nonlinear 1/ncore feature




Figure 48: Completion time vs ncores for query R2 with datasize 1000GB

3.3 Query R3

3.3.1 Query R3 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0780 | 0.9894 | 3915 | 0.0252 |
| Linear SVR | 0.1057 | 0.9816 | 5093 | 0.0299 |
| Polynomial SVR (2) | 0.7201 | 0.1170 | 41915 | 0.2811 |
| Polynomial SVR (3) | 0.2708 | 0.9232 | 15247 | 0.0945 |
| Polynomial SVR (4) | 0.6645 | 0.2415 | 36141 | 0.2452 |
| Polynomial SVR (6) | 0.6164 | 0.4095 | 34311 | 0.2194 |
| Gaussian SVR | 0.1788 | 0.9665 | 8642 | 0.0656 |

Table 49: Results for R3-250GB with the nonlinear $1/\text{ncores}$ feature




Figure 49: Completion time vs ncores for query R3 with datasize 250GB

3.3.2 Query R3 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0539 | 0.9962 | 10845 | 0.0220 |
| Linear SVR | 0.0626 | 0.9949 | 12632 | 0.0269 |
| Polynomial SVR (2) | 0.2821 | 0.8979 | 59067 | 0.1231 |
| Polynomial SVR (3) | 0.1297 | 0.9782 | 32168 | 0.0654 |
| Polynomial SVR (4) | 0.1922 | 0.9545 | 40999 | 0.0889 |
| Polynomial SVR (6) | 0.2877 | 0.8937 | 54259 | 0.1205 |
| Gaussian SVR | 0.0703 | 0.9939 | 15349 | 0.0321 |

Table 50: Results for R3-500GB with the nonlinear 1/ncores feature




Figure 50: Completion time vs ncores for query R3 with datasize 500GB

3.3.3 Query R3 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0209 | 0.9996 | 2658 | 0.0033 |
| Linear SVR | 0.0806 | 0.9961 | 11228 | 0.0149 |
| Polynomial SVR (2) | 1.0051 | 0.5158 | 140581 | 0.1721 |
| Polynomial SVR (3) | 0.4684 | 0.9239 | 65851 | 0.0849 |
| Polynomial SVR (4) | 0.7927 | 0.6494 | 111418 | 0.1502 |
| Polynomial SVR (6) | 0.6810 | 0.6845 | 95464 | 0.1247 |
| Gaussian SVR | 0.2061 | 0.9777 | 26875 | 0.0359 |

Table 51: Results for R3-750GB with the nonlinear 1/ncores feature




Figure 51: Completion time vs ncores for query R3 with datasize 750GB

3.3.4 Query R3 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0723 | 0.9956 | 10267 | 0.0094 |
| Linear SVR | 0.0959 | 0.9948 | 20283 | 0.0199 |
| Polynomial SVR (2) | 0.6954 | 0.6772 | 148312 | 0.1435 |
| Polynomial SVR (3) | 0.3570 | 0.9367 | 58398 | 0.0507 |
| Polynomial SVR (4) | 0.8566 | 0.5419 | 174113 | 0.1625 |
| Polynomial SVR (6) | 0.8709 | 0.4828 | 171350 | 0.1591 |
| Gaussian SVR | 0.2990 | 0.9557 | 49677 | 0.0438 |

Table 52: Results for R3-1000GB with the nonlinear 1/ncore feature




Figure 52: Completion time vs ncores for query R3 with datasize 1000GB

3.4 Query R4

3.4.1 Query R4 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.1407 | 0.9793 | 4859 | 0.0282 |
| Linear SVR | 0.1048 | 0.9893 | 4109 | 0.0239 |
| Polynomial SVR (2) | 0.8143 | 0.4103 | 33794 | 0.2429 |
| Polynomial SVR (3) | 0.4882 | 0.8326 | 18105 | 0.1119 |
| Polynomial SVR (4) | 0.8747 | 0.5418 | 37483 | 0.2594 |
| Polynomial SVR (6) | 0.8999 | 0.5869 | 38677 | 0.2634 |
| Gaussian SVR | 0.2798 | 0.9470 | 9282 | 0.0572 |

Table 53: Results for R4-250GB with the nonlinear 1/ncores feature




Figure 53: Completion time vs ncores for query R4 with datasize 250GB

3.4.2 Query R4 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0501 | 0.9946 | 5178 | 0.0145 |
| Linear SVR | 0.0850 | 0.9872 | 14922 | 0.0397 |
| Polynomial SVR (2) | 0.2128 | 0.9400 | 29712 | 0.0914 |
| Polynomial SVR (3) | 0.1287 | 0.9766 | 19858 | 0.0507 |
| Polynomial SVR (4) | 0.2801 | 0.9035 | 42883 | 0.1152 |
| Polynomial SVR (6) | 0.2924 | 0.8767 | 40499 | 0.1112 |
| Gaussian SVR | 0.0791 | 0.9876 | 11411 | 0.0287 |

Table 54: Results for R4-500GB with the nonlinear 1/ncores feature




Figure 54: Completion time vs ncores for query R4 with datasize 500GB

3.4.3 Query R4 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0216 | 0.9996 | 2333 | 0.0038 |
| Linear SVR | 0.0773 | 0.9968 | 8179 | 0.0146 |
| Polynomial SVR (2) | 0.6766 | 0.7235 | 82143 | 0.1361 |
| Polynomial SVR (3) | 0.3497 | 0.9401 | 37764 | 0.0591 |
| Polynomial SVR (4) | 0.7932 | 0.7572 | 91479 | 0.1430 |
| Polynomial SVR (6) | 0.9234 | 0.6666 | 101239 | 0.1550 |
| Gaussian SVR | 0.1553 | 0.9830 | 13013 | 0.0221 |

Table 55: Results for R4-750GB with the nonlinear 1/ncores feature




Figure 55: Completion time vs ncores for query R4 with datasize 750GB

3.4.4 Query R4 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0812 | 0.9944 | 45917 | 0.0244 |
| Linear SVR | 0.0785 | 0.9966 | 46627 | 0.0315 |
| Polynomial SVR (2) | 0.7046 | 0.6535 | 446048 | 0.2688 |
| Polynomial SVR (3) | 0.3197 | 0.9286 | 196869 | 0.1466 |
| Polynomial SVR (4) | 0.8131 | 0.7058 | 578071 | 0.4145 |
| Polynomial SVR (6) | 0.9574 | 0.5389 | 680038 | 0.4961 |
| Gaussian SVR | 0.2282 | 0.9582 | 129449 | 0.0703 |

Table 56: Results for R4-1000GB with the nonlinear 1/ncore feature




Figure 56: Completion time vs ncores for query R4 with datasize 1000GB

3.5 Query R5

3.5.1 Query R5 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 1.1046 | -0.0599 | 720 | 0.0277 |
| Linear SVR | 1.0254 | 0.3409 | 642 | 0.0247 |
| Polynomial SVR (2) | 1.1448 | 0.0293 | 773 | 0.0301 |
| Polynomial SVR (3) | 1.0585 | 0.1048 | 706 | 0.0273 |
| Polynomial SVR (4) | 1.0900 | 0.0650 | 729 | 0.0283 |
| Polynomial SVR (6) | 1.0613 | 0.0496 | 679 | 0.0264 |
| Gaussian SVR | 0.9708 | 0.3643 | 606 | 0.0234 |

Table 57: Results for R5-250GB with the nonlinear 1/ncores feature




Figure 57: Completion time vs ncores for query R5 with datasize 250GB

3.5.2 Query R5 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.4987 | 0.5323 | 389 | 0.0164 |
| Linear SVR | 0.4932 | 0.5986 | 381 | 0.0161 |
| Polynomial SVR (2) | 0.7350 | 0.2975 | 705 | 0.0297 |
| Polynomial SVR (3) | 0.6196 | 0.4052 | 521 | 0.0220 |
| Polynomial SVR (4) | 0.7239 | 0.2502 | 636 | 0.0265 |
| Polynomial SVR (6) | 0.6725 | 0.1906 | 683 | 0.0289 |
| Gaussian SVR | 0.5501 | 0.5643 | 416 | 0.0175 |

Table 58: Results for R5-500GB with the nonlinear 1/ncores feature




Figure 58: Completion time vs ncores for query R5 with datasize 500GB

3.5.3 Query R5 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.5758 | 0.6745 | 442 | 0.0177 |
| Linear SVR | 0.6121 | 0.7041 | 466 | 0.0186 |
| Polynomial SVR (2) | 2.5714 | 0.0675 | 1530 | 0.0618 |
| Polynomial SVR (3) | 0.8141 | 0.5165 | 830 | 0.0340 |
| Polynomial SVR (4) | 1.9622 | 0.0219 | 1418 | 0.0570 |
| Polynomial SVR (6) | 2.1235 | 0.0003 | 1574 | 0.0630 |
| Gaussian SVR | 0.6321 | 0.6870 | 486 | 0.0195 |

Table 59: Results for R5-750GB with the nonlinear 1/ncores feature




Figure 59: Completion time vs ncores for query R5 with datasize 750GB

3.5.4 Query R5 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.5047 | 0.1390 | 1333 | 0.0341 |
| Linear SVR | 0.4982 | 0.3487 | 1237 | 0.0316 |
| Polynomial SVR (2) | 0.4124 | 0.5666 | 873 | 0.0219 |
| Polynomial SVR (3) | 0.6250 | 0.2996 | 1597 | 0.0404 |
| Polynomial SVR (4) | 0.5453 | 0.3911 | 1321 | 0.0332 |
| Polynomial SVR (6) | 0.5542 | 0.2966 | 1421 | 0.0358 |
| Gaussian SVR | 0.3568 | 0.5856 | 851 | 0.0217 |

Table 60: Results for R5-1000GB with the nonlinear 1/ncore feature




Figure 60: Completion time vs ncores for query R5 with datasize 1000GB

4 Fixed Datasize, $ncores^{-1}$, only $ncores$

4.1 Query R1

4.1.1 Query R1 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 1.2438 | 0.3259 | 12279 | 0.1330 |
| Linear SVR | 1.2946 | 0.3792 | 11957 | 0.1258 |
| Polynomial SVR (2) | 1.5679 | 0.0009 | 21296 | 0.3295 |
| Polynomial SVR (3) | 1.3800 | 0.2700 | 13385 | 0.1465 |
| Polynomial SVR (4) | 1.5428 | 0.0283 | 19555 | 0.2835 |
| Polynomial SVR (6) | 1.4977 | 0.0704 | 17963 | 0.2575 |
| Gaussian SVR | 1.2913 | 0.4092 | 11896 | 0.1246 |

Table 61: Results for R1-250GB with the nonlinear $1/ncores$ feature, only $ncores$




Figure 61: Completion time vs $ncores$ for query R1 with datasize 250GB

4.1.2 Query R1 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.4321 | 0.8126 | 30237 | 0.2313 |
| Linear SVR | 0.4590 | 0.8228 | 29651 | 0.2051 |
| Polynomial SVR (2) | 0.7453 | 0.4527 | 56372 | 0.5019 |
| Polynomial SVR (3) | 0.3390 | 0.9148 | 25984 | 0.2401 |
| Polynomial SVR (4) | 0.5479 | 0.7103 | 41171 | 0.3775 |
| Polynomial SVR (6) | 0.4578 | 0.8009 | 34535 | 0.3294 |
| Gaussian SVR | 0.3841 | 0.9071 | 29393 | 0.2724 |

Table 62: Results for R1-500GB with the nonlinear 1/ncore feature, only ncores




Figure 62: Completion time vs ncores for query R1 with datasize 500GB

4.1.3 Query R1 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.5526 | 0.8091 | 42884 | 0.2200 |
| Linear SVR | 0.5520 | 0.8666 | 39858 | 0.2325 |
| Polynomial SVR (2) | 0.9850 | 0.6193 | 74917 | 0.4174 |
| Polynomial SVR (3) | 0.5397 | 0.8813 | 38123 | 0.2276 |
| Polynomial SVR (4) | 0.8382 | 0.7906 | 57443 | 0.3611 |
| Polynomial SVR (6) | 0.7219 | 0.8510 | 49401 | 0.3107 |
| Gaussian SVR | 0.5938 | 0.8893 | 42964 | 0.2585 |

Table 63: Results for R1-750GB with the nonlinear 1/ncores feature, only ncores




Figure 63: Completion time vs ncores for query R1 with datasize 750GB

4.1.4 Query R1 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.3229 | 0.8946 | 22902 | 0.0553 |
| Linear SVR | 0.3305 | 0.8955 | 22857 | 0.0562 |
| Polynomial SVR (2) | 0.8438 | 0.4005 | 71595 | 0.1618 |
| Polynomial SVR (3) | 0.3683 | 0.8635 | 22807 | 0.0558 |
| Polynomial SVR (4) | 0.6342 | 0.6012 | 44393 | 0.1110 |
| Polynomial SVR (6) | 0.5642 | 0.6945 | 37885 | 0.0928 |
| Gaussian SVR | 0.3441 | 0.9190 | 26107 | 0.0602 |

Table 64: Results for R1-1000GB with the nonlinear 1/ncores feature, only ncores




Figure 64: Completion time vs ncores for query R1 with datasize 1000GB

4.2 Query R2

4.2.1 Query R2 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.9065 | -0.0568 | 2785 | 0.0339 |
| Linear SVR | 0.9191 | 0.0883 | 2817 | 0.0343 |
| Polynomial SVR (2) | 0.9837 | 0.0002 | 3054 | 0.0373 |
| Polynomial SVR (3) | 0.9247 | 0.0715 | 2809 | 0.0342 |
| Polynomial SVR (4) | 0.9515 | 0.0076 | 2885 | 0.0351 |
| Polynomial SVR (6) | 0.9422 | 0.0161 | 2813 | 0.0342 |
| Gaussian SVR | 0.9145 | 0.1079 | 2823 | 0.0344 |

Table 65: Results for R2-250GB with the nonlinear $1/\text{ncores}$ feature, only ncores




Figure 65: Completion time vs ncores for query R2 with datasize 250GB

4.2.2 Query R2 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 1.2224 | -0.0227 | 2953 | 0.0406 |
| Linear SVR | 1.2263 | 0.0000 | 2982 | 0.0411 |
| Polynomial SVR (2) | 1.1928 | 0.1744 | 2978 | 0.0411 |
| Polynomial SVR (3) | 1.2252 | 0.0075 | 2985 | 0.0411 |
| Polynomial SVR (4) | 1.2031 | 0.1128 | 2964 | 0.0408 |
| Polynomial SVR (6) | 1.2064 | 0.0870 | 2956 | 0.0407 |
| Gaussian SVR | 1.2449 | 0.0043 | 3126 | 0.0432 |

Table 66: Results for R2-500GB with the nonlinear $1/\text{ncores}$ feature, only ncores




Figure 66: Completion time vs ncores for query R2 with datasize 500GB

4.2.3 Query R2 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 1.0119 | -0.0671 | 2082 | 0.0266 |
| Linear SVR | 1.0201 | 0.0504 | 2009 | 0.0255 |
| Polynomial SVR (2) | 1.0331 | 0.0354 | 2078 | 0.0264 |
| Polynomial SVR (3) | 1.0269 | 0.0655 | 2050 | 0.0261 |
| Polynomial SVR (4) | 1.0306 | 0.0485 | 2052 | 0.0261 |
| Polynomial SVR (6) | 1.0210 | 0.0548 | 2045 | 0.0260 |
| Gaussian SVR | 1.0305 | 0.0202 | 2063 | 0.0262 |

Table 67: Results for R2-750GB with the nonlinear $1/\text{ncores}$ feature, only ncores




Figure 67: Completion time vs ncores for query R2 with datasize 750GB

4.2.4 Query R2 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.3258 | 0.8006 | 157650 | 0.1926 |
| Linear SVR | 0.4842 | 0.8331 | 223250 | 0.2752 |
| Polynomial SVR (2) | 0.6883 | 0.1758 | 329819 | 0.3776 |
| Polynomial SVR (3) | 0.4559 | 0.7612 | 173610 | 0.2351 |
| Polynomial SVR (4) | 0.5423 | 0.4599 | 258020 | 0.2906 |
| Polynomial SVR (6) | 0.4626 | 0.6044 | 229115 | 0.2737 |
| Gaussian SVR | 0.1730 | 0.9691 | 69502 | 0.0818 |

Table 68: Results for R2-1000GB with the nonlinear 1/ncores feature, only ncores




Figure 68: Completion time vs ncores for query R2 with datasize 1000GB

4.3 Query R3

4.3.1 Query R3 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.7627 | 0.4444 | 40061 | 0.1892 |
| Linear SVR | 0.7863 | 0.7014 | 40887 | 0.1865 |
| Polynomial SVR (2) | 1.0935 | 0.1924 | 61331 | 0.3228 |
| Polynomial SVR (3) | 0.8169 | 0.6371 | 43254 | 0.2055 |
| Polynomial SVR (4) | 0.9730 | 0.3470 | 54074 | 0.2851 |
| Polynomial SVR (6) | 1.0190 | 0.4264 | 57989 | 0.2907 |
| Gaussian SVR | 0.7857 | 0.6768 | 42184 | 0.2023 |

Table 69: Results for R3-250GB with the nonlinear $1/\text{ncores}$ feature, only ncores




Figure 69: Completion time vs ncores for query R3 with datasize 250GB

4.3.2 Query R3 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.3162 | 0.8991 | 74328 | 0.1622 |
| Linear SVR | 0.3102 | 0.9030 | 71562 | 0.1475 |
| Polynomial SVR (2) | 0.8017 | 0.3717 | 170794 | 0.4134 |
| Polynomial SVR (3) | 0.1738 | 0.9698 | 33539 | 0.0823 |
| Polynomial SVR (4) | 0.6309 | 0.6249 | 122119 | 0.3179 |
| Polynomial SVR (6) | 0.5001 | 0.7625 | 102471 | 0.2574 |
| Gaussian SVR | 0.1390 | 0.9814 | 27308 | 0.0611 |

Table 70: Results for R3-500GB with the nonlinear 1/ncores feature, only ncores




Figure 70: Completion time vs ncores for query R3 with datasize 500GB

4.3.3 Query R3 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.2049 | 0.9524 | 28016 | 0.0404 |
| Linear SVR | 0.2057 | 0.9575 | 28111 | 0.0401 |
| Polynomial SVR (2) | 0.8224 | 0.2464 | 128039 | 0.1753 |
| Polynomial SVR (3) | 0.3301 | 0.8941 | 46655 | 0.0635 |
| Polynomial SVR (4) | 0.6989 | 0.5009 | 93688 | 0.1338 |
| Polynomial SVR (6) | 0.5836 | 0.6402 | 78143 | 0.1099 |
| Gaussian SVR | 0.1629 | 0.9849 | 18453 | 0.0281 |

Table 71: Results for R3-750GB with the nonlinear 1/ncore feature, only ncores




Figure 71: Completion time vs ncores for query R3 with datasize 750GB

4.3.4 Query R3 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.2395 | 0.9391 | 47342 | 0.0519 |
| Linear SVR | 0.2429 | 0.9424 | 47915 | 0.0527 |
| Polynomial SVR (2) | 0.7601 | 0.4483 | 150763 | 0.1636 |
| Polynomial SVR (3) | 0.3880 | 0.8856 | 63063 | 0.0725 |
| Polynomial SVR (4) | 0.5732 | 0.6936 | 111389 | 0.1218 |
| Polynomial SVR (6) | 0.4757 | 0.7969 | 94384 | 0.1027 |
| Gaussian SVR | 0.0893 | 0.9941 | 15506 | 0.0164 |

Table 72: Results for R3-1000GB with the nonlinear $1/\text{ncores}$ feature, only ncores




Figure 72: Completion time vs ncores for query R3 with datasize 1000GB

4.4 Query R4

4.4.1 Query R4 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.8686 | 0.1477 | 36477 | 0.2638 |
| Linear SVR | 0.8530 | 0.2189 | 36318 | 0.2718 |
| Polynomial SVR (2) | 1.0554 | 0.0583 | 42311 | 0.3487 |
| Polynomial SVR (3) | 0.9447 | 0.0901 | 37389 | 0.2771 |
| Polynomial SVR (4) | 1.0542 | 0.0017 | 43480 | 0.3472 |
| Polynomial SVR (6) | 1.0304 | 0.0023 | 42378 | 0.3295 |
| Gaussian SVR | 0.8335 | 0.2617 | 35373 | 0.2615 |

Table 73: Results for R4-250GB with the nonlinear $1/\text{ncores}$ feature, only ncores




Figure 73: Completion time vs ncores for query R4 with datasize 250GB

4.4.2 Query R4 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.3338 | 0.9217 | 56936 | 0.1305 |
| Linear SVR | 0.3256 | 0.9357 | 55121 | 0.1264 |
| Polynomial SVR (2) | 0.9065 | 0.4895 | 144214 | 0.3545 |
| Polynomial SVR (3) | 0.2259 | 0.9863 | 34171 | 0.0723 |
| Polynomial SVR (4) | 0.6572 | 0.7210 | 94909 | 0.2849 |
| Polynomial SVR (6) | 0.5188 | 0.8340 | 78746 | 0.2291 |
| Gaussian SVR | 0.0987 | 0.9941 | 15135 | 0.0407 |

Table 74: Results for R4-500GB with the nonlinear 1/ncore feature, only ncores




Figure 74: Completion time vs ncores for query R4 with datasize 500GB

4.4.3 Query R4 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.2516 | 0.8972 | 26818 | 0.0505 |
| Linear SVR | 0.2485 | 0.9195 | 27232 | 0.0506 |
| Polynomial SVR (2) | 0.6507 | 0.3168 | 79129 | 0.1386 |
| Polynomial SVR (3) | 0.4005 | 0.7802 | 41995 | 0.0791 |
| Polynomial SVR (4) | 0.5367 | 0.5595 | 58504 | 0.1007 |
| Polynomial SVR (6) | 0.4790 | 0.6565 | 54216 | 0.0925 |
| Gaussian SVR | 0.1194 | 0.9780 | 12187 | 0.0212 |

Table 75: Results for R4-750GB with the nonlinear $1/\text{ncores}$ feature, only ncores




Figure 75: Completion time vs ncores for query R4 with datasize 750GB

4.4.4 Query R4 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.6309 | 0.3868 | 402429 | 0.3275 |
| Linear SVR | 0.6552 | 0.4433 | 491316 | 0.3635 |
| Polynomial SVR (2) | 0.8742 | 0.0220 | 656146 | 0.4904 |
| Polynomial SVR (3) | 0.7401 | 0.2208 | 516331 | 0.3896 |
| Polynomial SVR (4) | 0.9535 | 0.0796 | 721224 | 0.5477 |
| Polynomial SVR (6) | 0.9594 | 0.1062 | 724319 | 0.5523 |
| Gaussian SVR | 0.2640 | 0.8931 | 147333 | 0.0879 |

Table 76: Results for R4-1000GB with the nonlinear $1/\text{ncores}$ feature, only ncores




Figure 76: Completion time vs ncores for query R4 with datasize 1000GB

4.5 Query R5

4.5.1 Query R5 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.8453 | -0.3635 | 534 | 0.0213 |
| Linear SVR | 0.7863 | 0.0052 | 514 | 0.0204 |
| Polynomial SVR (2) | 0.7718 | 0.0620 | 498 | 0.0198 |
| Polynomial SVR (3) | 0.7844 | 0.0178 | 515 | 0.0204 |
| Polynomial SVR (4) | 0.8009 | 0.0468 | 502 | 0.0199 |
| Polynomial SVR (6) | 0.8011 | 0.0419 | 497 | 0.0197 |
| Gaussian SVR | 0.8189 | 0.0012 | 486 | 0.0192 |

Table 77: Results for R5-250GB with the nonlinear $1/\text{ncores}$ feature, only ncores




Figure 77: Completion time vs ncores for query R5 with datasize 250GB

4.5.2 Query R5 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 1.1189 | -0.0289 | 1200 | 0.0483 |
| Linear SVR | 1.1011 | 0.1510 | 1214 | 0.0493 |
| Polynomial SVR (2) | 1.1335 | 0.0115 | 1295 | 0.0528 |
| Polynomial SVR (3) | 1.0972 | 0.1235 | 1221 | 0.0497 |
| Polynomial SVR (4) | 1.1256 | 0.0000 | 1275 | 0.0519 |
| Polynomial SVR (6) | 1.1101 | 0.0053 | 1273 | 0.0520 |
| Gaussian SVR | 1.1255 | 0.0005 | 1276 | 0.0520 |

Table 78: Results for R5-500GB with the nonlinear $1/\text{ncores}$ feature, only ncores




Figure 78: Completion time vs ncores for query R5 with datasize 500GB

4.5.3 Query R5 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.8980 | -0.0354 | 1068 | 0.0445 |
| Linear SVR | 1.0199 | -0.0000 | 983 | 0.0425 |
| Polynomial SVR (2) | 1.0199 | -0.0000 | 983 | 0.0425 |
| Polynomial SVR (3) | 1.0199 | -0.0000 | 983 | 0.0425 |
| Polynomial SVR (4) | 1.0199 | -0.0000 | 983 | 0.0425 |
| Polynomial SVR (6) | 1.0199 | -0.0000 | 983 | 0.0425 |
| Gaussian SVR | 1.0199 | -0.0000 | 983 | 0.0425 |

Table 79: Results for R5-750GB with the nonlinear 1/ncore feature, only ncores




Figure 79: Completion time vs ncores for query R5 with datasize 750GB

4.5.4 Query R5 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.9881 | 0.2649 | 3615 | 0.0851 |
| Linear SVR | 1.0305 | 0.2927 | 3493 | 0.0805 |
| Polynomial SVR (2) | 1.0317 | 0.2871 | 3100 | 0.0694 |
| Polynomial SVR (3) | 1.0104 | 0.3333 | 3346 | 0.0763 |
| Polynomial SVR (4) | 1.0033 | 0.3402 | 2998 | 0.0667 |
| Polynomial SVR (6) | 1.0049 | 0.3490 | 3055 | 0.0682 |
| Gaussian SVR | 0.9937 | 0.3578 | 3236 | 0.0737 |

Table 80: Results for R5-1000GB with the nonlinear $1/\text{ncores}$ feature, only ncores




Figure 80: Completion time vs ncores for query R5 with datasize 1000GB

5 Query comparison

5.1 R5 \rightarrow R2 (Fast - Fast)

5.1.1 R5 \rightarrow R2 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.2664 | -3.8012 | 7131 | 0.0868 |
| Linear SVR | 1.4192 | 0.4500 | 41150 | 0.4928 |
| Polynomial SVR (2) | 2.0017 | 0.3318 | 57963 | 0.6936 |
| Polynomial SVR (3) | 1.7351 | 0.3201 | 50336 | 0.6032 |
| Polynomial SVR (4) | 2.0052 | 0.2591 | 58051 | 0.6946 |
| Polynomial SVR (6) | 2.0177 | 0.3148 | 58314 | 0.6974 |
| Gaussian SVR | 1.9531 | 0.0953 | 56646 | 0.6783 |

Table 81: Results for R5 \rightarrow R2-250




Figure 81: Completion time vs ncores for: R5 \rightarrow R2 with datasize 250GB

5.1.2 R5 \rightarrow R2 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0927 | 0.4825 | 1627 | 0.0219 |
| Linear SVR | 0.8177 | 0.8804 | 20156 | 0.2760 |
| Polynomial SVR (2) | 1.9565 | 0.8888 | 47958 | 0.6544 |
| Polynomial SVR (3) | 1.3110 | 0.8506 | 32346 | 0.4437 |
| Polynomial SVR (4) | 1.9813 | 0.8130 | 48399 | 0.6598 |
| Polynomial SVR (6) | 1.9723 | 0.7576 | 48394 | 0.6606 |
| Gaussian SVR | 1.8458 | 0.4395 | 45448 | 0.6213 |

Table 82: Results for R5 \rightarrow R2-500




Figure 82: Completion time vs ncores for: R5 \rightarrow R2 with datasize 500GB

5.1.3 R5 \rightarrow R2 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.7284 | - | 19886 | 0.2523 |
| Linear SVR | 1.2200 | 0.7629 | 33371 | 0.4226 |
| Polynomial SVR (2) | 1.9610 | 0.7081 | 53638 | 0.6792 |
| Polynomial SVR (3) | 1.7155 | 0.7155 | 46976 | 0.5953 |
| Polynomial SVR (4) | 1.9544 | 0.6962 | 53405 | 0.6760 |
| Polynomial SVR (6) | 1.9545 | 0.6360 | 53347 | 0.6749 |
| Gaussian SVR | 1.8982 | 0.4132 | 51944 | 0.6579 |

Table 83: Results for R5 \rightarrow R2-750




Figure 83: Completion time vs ncores for: R5 \rightarrow R2 with datasize 750GB

5.1.4 R5 \rightarrow R2 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.2791 | 0.8182 | 124323 | 0.1468 |
| Linear SVR | 1.9131 | -0.0000 | 865805 | 0.9413 |
| Polynomial SVR (2) | 1.9131 | -0.0000 | 865805 | 0.9413 |
| Polynomial SVR (3) | 1.9131 | -0.0000 | 865805 | 0.9413 |
| Polynomial SVR (4) | 1.9131 | -0.0000 | 865805 | 0.9413 |
| Polynomial SVR (6) | 1.9131 | -0.0000 | 865805 | 0.9413 |
| Gaussian SVR | 1.9131 | -0.0000 | 865805 | 0.9413 |

Table 84: Results for R5 \rightarrow R2-1000




Figure 84: Completion time vs ncores for: R5 \rightarrow R2 with datasize 1000GB

5.2 R2 → R5 (Fast - Fast)

5.2.1 R2 → R5 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.5076 | - 54.2010 | 14618 | 0.5690 |
| Linear SVR | 0.0742 | 0.5354 | 1906 | 0.0748 |
| Polynomial SVR (2) | 1.9582 | 0.5005 | 56793 | 2.2284 |
| Polynomial SVR (3) | 1.4334 | 0.4575 | 41627 | 1.6299 |
| Polynomial SVR (4) | 1.9578 | 0.4334 | 56818 | 2.2282 |
| Polynomial SVR (6) | 1.9529 | 0.3733 | 56692 | 2.2226 |
| Gaussian SVR | 1.8646 | 0.4311 | 54134 | 2.1221 |

Table 85: Results for R2 → R5-250




Figure 85: Completion time vs ncores for: R2 → R5 with datasize 250GB

5.2.2 R2 → R5 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.2910 | -3.7899 | 6672 | 0.2718 |
| Linear SVR | 0.7805 | 0.3478 | 19144 | 0.8083 |
| Polynomial SVR (2) | 1.9764 | 0.3328 | 48648 | 2.0492 |
| Polynomial SVR (3) | 1.5362 | 0.3112 | 37827 | 1.5916 |
| Polynomial SVR (4) | 1.9546 | 0.2841 | 48117 | 2.0265 |
| Polynomial SVR (6) | 1.9390 | 0.2654 | 47733 | 2.0103 |
| Gaussian SVR | 1.9761 | 0.1983 | 48681 | 2.0483 |

Table 86: Results for R2 → R5-500




Figure 86: Completion time vs ncores for: R2 → R5 with datasize 500GB

5.2.3 R2 \rightarrow R5 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.7050 | - | 19273 | 0.7906 |
| Linear SVR | 0.5814 | 0.1630 | 15781 | 0.6549 |
| Polynomial SVR (2) | 1.9918 | 0.3443 | 54498 | 2.2536 |
| Polynomial SVR (3) | 1.5014 | 0.3534 | 41121 | 1.6974 |
| Polynomial SVR (4) | 1.9754 | 0.2597 | 54059 | 2.2349 |
| Polynomial SVR (6) | 1.9710 | 0.2735 | 53943 | 2.2300 |
| Gaussian SVR | 2.0285 | 0.4331 | 55559 | 2.2948 |

Table 87: Results for R2 \rightarrow R5-750




Figure 87: Completion time vs ncores for: R2 \rightarrow R5 with datasize 750GB

5.2.4 R2 \rightarrow R5 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.2457 | - | 117597 | 2.9684 |
| Linear SVR | 0.0746 | 0.0791 | 31327 | 0.7786 |
| Polynomial SVR (2) | 1.5254 | 0.2338 | 734175 | 18.4731 |
| Polynomial SVR (3) | 0.8871 | 0.1001 | 426009 | 10.6683 |
| Polynomial SVR (4) | 1.4660 | 0.0268 | 705234 | 17.7261 |
| Polynomial SVR (6) | 1.4552 | 0.0225 | 700727 | 17.5862 |
| Gaussian SVR | 2.1204 | 0.0853 | 1020795 | 25.5869 |

Table 88: Results for R2 \rightarrow R5-1000




Figure 88: Completion time vs ncores for: R2 \rightarrow R5 with datasize 1000GB

5.3 R3 → R4 (Slow - Slow)

5.3.1 R3 → R4 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0999 | 0.9905 | 7743 | 0.0556 |
| Linear SVR | 0.1017 | 0.9936 | 7894 | 0.0559 |
| Polynomial SVR (2) | 1.2864 | 0.0000 | 109952 | 0.8321 |
| Polynomial SVR (3) | 0.5276 | 0.8127 | 40426 | 0.2156 |
| Polynomial SVR (4) | 0.9298 | 0.2785 | 77573 | 0.5374 |
| Polynomial SVR (6) | 0.8746 | 0.2982 | 69520 | 0.4722 |
| Gaussian SVR | 0.5125 | 0.8738 | 39988 | 0.2815 |

Table 89: Results for R3 → R4-250




Figure 89: Completion time vs ncores for: R3 → R4 with datasize 250GB

5.3.2 R3 → R4 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.5575 | 0.7917 | 79305 | 0.1806 |
| Linear SVR | 0.6248 | 0.8895 | 87857 | 0.1990 |
| Polynomial SVR (2) | 1.4064 | 0.3856 | 175966 | 0.4039 |
| Polynomial SVR (3) | 0.9288 | 0.7713 | 129311 | 0.3215 |
| Polynomial SVR (4) | 1.3185 | 0.1435 | 177714 | 0.4267 |
| Polynomial SVR (6) | 1.4713 | 0.1476 | 177441 | 0.4040 |
| Gaussian SVR | 1.0797 | 0.3126 | 134221 | 0.2887 |

Table 90: Results for R3 → R4-500




Figure 90: Completion time vs ncores for: R3 → R4 with datasize 500GB

5.3.3 R3 → R4 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0933 | 0.9899 | 12026 | 0.0195 |
| Linear SVR | 0.2629 | 0.9912 | 33564 | 0.0561 |
| Polynomial SVR (2) | 1.8337 | 0.7689 | 217984 | 0.3828 |
| Polynomial SVR (3) | 0.7352 | 0.7039 | 87799 | 0.1548 |
| Polynomial SVR (4) | 1.8539 | 0.4641 | 210264 | 0.3621 |
| Polynomial SVR (6) | 1.1873 | 0.4401 | 148303 | 0.2645 |
| Gaussian SVR | 1.0774 | 0.8421 | 126509 | 0.2349 |

Table 91: Results for R3 → R4-750




Figure 91: Completion time vs ncores for: R3 → R4 with datasize 750GB

5.3.4 R3 → R4 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 1.7878 | -1.2231 | 784437 | 0.3797 |
| Linear SVR | 1.6361 | 0.0281 | 660983 | 0.2902 |
| Polynomial SVR (2) | 1.6549 | 0.0157 | 689481 | 0.3172 |
| Polynomial SVR (3) | 1.6757 | 0.0650 | 705047 | 0.3218 |
| Polynomial SVR (4) | 1.9237 | 0.1209 | 764929 | 0.3431 |
| Polynomial SVR (6) | 1.8627 | 0.0254 | 735220 | 0.3294 |
| Gaussian SVR | 1.5644 | 0.0772 | 640303 | 0.2875 |

Table 92: Results for R3 → R4-1000




Figure 92: Completion time vs ncores for: R3 → R4 with datasize 1000GB

5.4 R4 → R3 (Slow - Slow)

5.4.1 R4 → R3 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.1117 | 0.9859 | 6402 | 0.0272 |
| Linear SVR | 0.1241 | 0.9882 | 8184 | 0.0379 |
| Polynomial SVR (2) | 1.0187 | 0.0102 | 84602 | 0.4395 |
| Polynomial SVR (3) | 0.7609 | 0.7115 | 53158 | 0.2639 |
| Polynomial SVR (4) | 1.0989 | 0.0178 | 85066 | 0.3771 |
| Polynomial SVR (6) | 1.0632 | 0.0244 | 82125 | 0.3725 |
| Gaussian SVR | 0.5420 | 0.8101 | 41963 | 0.2180 |

Table 93: Results for R4 → R3-250




Figure 93: Completion time vs ncores for: R4 → R3 with datasize 250GB

5.4.2 R4 \rightarrow R3 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.1181 | 0.9592 | 13349 | 0.0297 |
| Linear SVR | 0.1281 | 0.9719 | 13961 | 0.0304 |
| Polynomial SVR (2) | 1.0739 | 0.0989 | 142216 | 0.3246 |
| Polynomial SVR (3) | 0.5917 | 0.7504 | 77226 | 0.1606 |
| Polynomial SVR (4) | 1.0211 | 0.2288 | 134494 | 0.3151 |
| Polynomial SVR (6) | 0.9462 | 0.0853 | 117567 | 0.2718 |
| Gaussian SVR | 0.3507 | 0.7951 | 46118 | 0.1032 |

Table 94: Results for R4 \rightarrow R3-500




Figure 94: Completion time vs ncores for: R4 \rightarrow R3 with datasize 500GB

5.4.3 R4 → R3 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.1065 | 0.9853 | 13983 | 0.0192 |
| Linear SVR | 0.1572 | 0.9937 | 18494 | 0.0250 |
| Polynomial SVR (2) | 1.7890 | 0.8077 | 185642 | 0.2313 |
| Polynomial SVR (3) | 1.0593 | 0.8125 | 126124 | 0.1704 |
| Polynomial SVR (4) | 2.8642 | 0.6965 | 233765 | 0.2775 |
| Polynomial SVR (6) | 2.9392 | 0.5268 | 213611 | 0.2468 |
| Gaussian SVR | 0.8538 | 0.7386 | 85514 | 0.1055 |

Table 95: Results for R4 → R3-750




Figure 95: Completion time vs ncores for: R4 → R3 with datasize 750GB

5.4.4 R4 → R3 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.2979 | 0.2999 | 163252 | 0.1736 |
| Linear SVR | 0.2150 | 0.9597 | 109503 | 0.1219 |
| Polynomial SVR (2) | 0.7033 | 0.6852 | 382105 | 0.4195 |
| Polynomial SVR (3) | 0.5805 | 0.7986 | 324151 | 0.3440 |
| Polynomial SVR (4) | 0.3638 | 0.7740 | 189564 | 0.2076 |
| Polynomial SVR (6) | 0.4766 | 0.6324 | 252222 | 0.2741 |
| Gaussian SVR | 1.0142 | 0.5173 | 574326 | 0.6115 |

Table 96: Results for R4 → R3-1000




Figure 96: Completion time vs ncores for: R4 → R3 with datasize 1000GB

5.5 R1, R2, R4 → R3 (Fast and Slow - Average)

5.5.1 R1,R2,R4 → R3 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.1484 | 0.9772 | 9501 | 0.0466 |
| Linear SVR | 0.1560 | 0.9845 | 10574 | 0.0553 |
| Polynomial SVR (2) | 0.8548 | 0.7060 | 57023 | 0.2555 |
| Polynomial SVR (3) | 0.4469 | 0.8616 | 28093 | 0.1225 |
| Polynomial SVR (4) | 0.7404 | 0.7503 | 42222 | 0.1783 |
| Polynomial SVR (6) | 1.1648 | 0.5504 | 61124 | 0.2455 |
| Gaussian SVR | 0.3926 | 0.9405 | 24779 | 0.1084 |

Table 97: Results for R1,R2,R4 → R3-250




Figure 97: Completion time vs ncores for: R1,R2,R4 → R3 with datasize 250GB

5.5.2 R1,R2,R4 \rightarrow R3 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0914 | 0.9687 | 13457 | 0.0317 |
| Linear SVR | 0.0833 | 0.9741 | 13357 | 0.0308 |
| Polynomial SVR (2) | 0.6833 | 0.6864 | 118761 | 0.2933 |
| Polynomial SVR (3) | 0.2660 | 0.8449 | 42425 | 0.0947 |
| Polynomial SVR (4) | 0.7291 | 0.6105 | 129234 | 0.3055 |
| Polynomial SVR (6) | 0.7959 | 0.4006 | 136789 | 0.3104 |
| Gaussian SVR | 0.2844 | 0.8990 | 43387 | 0.1147 |

Table 98: Results for R1,R2,R4 \rightarrow R3-500




Figure 98: Completion time vs ncores for: R1,R2,R4 \rightarrow R3 with datasize 500GB

5.5.3 R1,R2,R4 \rightarrow R3 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0581 | 0.9694 | 13358 | 0.0186 |
| Linear SVR | 0.0432 | 0.9924 | 10296 | 0.0152 |
| Polynomial SVR (2) | 0.4219 | 0.7036 | 106082 | 0.1502 |
| Polynomial SVR (3) | 0.2763 | 0.7323 | 53615 | 0.0742 |
| Polynomial SVR (4) | 0.6688 | 0.4095 | 168830 | 0.2349 |
| Polynomial SVR (6) | 0.6903 | 0.2073 | 166482 | 0.2269 |
| Gaussian SVR | 0.3779 | 0.5523 | 82501 | 0.1133 |

Table 99: Results for R1,R2,R4 \rightarrow R3-750




Figure 99: Completion time vs ncores for: R1,R2,R4 \rightarrow R3 with datasize 750GB

5.5.4 R1,R2,R4 \rightarrow R3 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.5423 | -0.7126 | 260530 | 0.2553 |
| Linear SVR | 0.5775 | 0.9500 | 275655 | 0.2718 |
| Polynomial SVR (2) | 0.5045 | 0.7467 | 202083 | 0.1809 |
| Polynomial SVR (3) | 0.4542 | 0.0345 | 148622 | 0.1290 |
| Polynomial SVR (4) | 0.5079 | 0.0024 | 181230 | 0.1545 |
| Polynomial SVR (6) | 0.5769 | 0.0175 | 203255 | 0.1711 |
| Gaussian SVR | 0.3920 | 0.7004 | 168171 | 0.1877 |

Table 100: Results for R1,R2,R4 \rightarrow R3-1000




Figure 100: Completion time vs ncores for: R1,R2,R4 \rightarrow R3 with datasize 1000GB

5.6 R2, R5 \rightarrow R3 (Fast - Slow)

5.6.1 R2,R5 \rightarrow R3 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 2.9358 | -8.9645 | 225067 | 1.0067 |
| Linear SVR | 2.0189 | 0.1000 | 154238 | 0.6860 |
| Polynomial SVR (2) | 2.3413 | 0.0143 | 179635 | 0.8056 |
| Polynomial SVR (3) | 2.0566 | 0.0575 | 157367 | 0.6993 |
| Polynomial SVR (4) | 2.1957 | 0.0474 | 167202 | 0.7392 |
| Polynomial SVR (6) | 2.1578 | 0.0358 | 163842 | 0.7215 |
| Gaussian SVR | 2.0866 | 0.2880 | 160725 | 0.7205 |

Table 101: Results for R2,R5 \rightarrow R3-250




Figure 101: Completion time vs ncores for: R2,R5 \rightarrow R3 with datasize 250GB

5.6.2 R2,R5 \rightarrow R3 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|---------|----------------|---------------------|---------------------|
| Linear regression | 24.4317 | - | 4082755 | 9.5054 |
| Linear SVR | 1.3901 | 0.0020 | 217391 | 0.4766 |
| Polynomial SVR (2) | 2.3279 | 0.8694 | 385016 | 0.8855 |
| Polynomial SVR (3) | 2.1472 | 0.6028 | 358017 | 0.8305 |
| Polynomial SVR (4) | 2.2733 | 0.4990 | 377968 | 0.8744 |
| Polynomial SVR (6) | 2.2643 | 0.2486 | 376874 | 0.8730 |
| Gaussian SVR | 2.2414 | 0.1919 | 373980 | 0.8683 |

Table 102: Results for R2,R5 \rightarrow R3-500




Figure 102: Completion time vs ncores for: R2,R5 \rightarrow R3 with datasize 500GB

5.6.3 R2,R5 \rightarrow R3 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 3.1762 | - | 905865 | 1.2847 |
| Linear SVR | 2.2593 | 0.0458 | 645146 | 0.9158 |
| Polynomial SVR (2) | 2.2682 | 0.1040 | 647657 | 0.9193 |
| Polynomial SVR (3) | 2.2586 | 0.3000 | 644988 | 0.9156 |
| Polynomial SVR (4) | 2.2552 | 0.0930 | 643893 | 0.9139 |
| Polynomial SVR (6) | 2.2549 | 0.0726 | 643840 | 0.9139 |
| Gaussian SVR | 2.2664 | 0.6737 | 647187 | 0.9187 |

Table 103: Results for R2,R5 \rightarrow R3-750




Figure 103: Completion time vs ncores for: R2,R5 \rightarrow R3 with datasize 750GB

5.6.4 R2,R5 \rightarrow R3 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.4825 | -0.2162 | 225269 | 0.2231 |
| Linear SVR | 1.5082 | 0.8905 | 700319 | 0.6897 |
| Polynomial SVR (2) | 1.8532 | 0.0474 | 856486 | 0.8417 |
| Polynomial SVR (3) | 1.7507 | 0.0017 | 810643 | 0.7966 |
| Polynomial SVR (4) | 1.8048 | 0.1072 | 838182 | 0.8256 |
| Polynomial SVR (6) | 1.8737 | 0.3832 | 871983 | 0.8603 |
| Gaussian SVR | 1.5367 | 0.8611 | 723592 | 0.7227 |

Table 104: Results for R2,R5 \rightarrow R3-1000




Figure 104: Completion time vs ncores for: R2,R5 \rightarrow R3 with datasize 1000GB

5.7 R3, R4 \rightarrow R2 (Fast - Slow)

5.7.1 R3,R4 \rightarrow R2 — Datasize 250GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.7320 | - | 67111 | 0.8046 |
| Linear SVR | 0.6184 | 0.3343 | 56625 | 0.6788 |
| Polynomial SVR (2) | 2.2958 | 0.1994 | 209908 | 2.5228 |
| Polynomial SVR (3) | 0.6067 | 0.0155 | 47847 | 0.5734 |
| Polynomial SVR (4) | 2.0229 | 0.0101 | 180533 | 2.1680 |
| Polynomial SVR (6) | 1.7103 | 0.0153 | 146100 | 1.7568 |
| Gaussian SVR | 1.0720 | 0.2857 | 98273 | 1.1766 |

Table 105: Results for R3,R4 \rightarrow R2-250




Figure 105: Completion time vs ncores for: R3,R4 \rightarrow R2 with datasize 250GB

5.7.2 R3,R4 \rightarrow R2 — Datasize 500GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.2274 | - | 50511 | 0.6952 |
| Linear SVR | 0.4418 | 0.0006 | 96855 | 1.3309 |
| Polynomial SVR (2) | 1.8749 | 0.0345 | 420366 | 5.7616 |
| Polynomial SVR (3) | 0.9866 | 0.0290 | 213001 | 2.9014 |
| Polynomial SVR (4) | 1.8453 | 0.0260 | 410413 | 5.6274 |
| Polynomial SVR (6) | 1.7365 | 0.0248 | 385605 | 5.2880 |
| Gaussian SVR | 1.4964 | 0.0442 | 337994 | 4.6265 |

Table 106: Results for R3,R4 \rightarrow R2-500




Figure 106: Completion time vs ncores for: R3,R4 \rightarrow R2 with datasize 500GB

5.7.3 R3,R4 \rightarrow R2 — Datasize 750GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.4540 | - | 146163 | 1.8730 |
| Linear SVR | 0.4483 | 0.0421 | 137083 | 1.7535 |
| Polynomial SVR (2) | 2.4760 | 0.1503 | 788379 | 10.1194 |
| Polynomial SVR (3) | 1.1909 | 0.3274 | 304007 | 3.8427 |
| Polynomial SVR (4) | 2.6951 | 0.2837 | 803028 | 10.3423 |
| Polynomial SVR (6) | 3.3215 | 0.2411 | 930987 | 11.9433 |
| Gaussian SVR | 1.9111 | 0.1845 | 615634 | 7.8865 |

Table 107: Results for R3,R4 \rightarrow R2-750




Figure 107: Completion time vs ncores for: R3,R4 \rightarrow R2 with datasize 750GB

5.7.4 R3,R4 \rightarrow R2 — Datasize 1000GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|---------|----------------|---------------------|---------------------|
| Linear regression | 11.2761 | - | 3410161 | 3.2357 |
| Linear SVR | 7.7162 | 0.9822 | 3657876 | 4.2191 |
| Polynomial SVR (2) | 8.0373 | 0.9441 | 2766694 | 2.6339 |
| Polynomial SVR (3) | 10.5398 | 0.8977 | 3750682 | 3.6335 |
| Polynomial SVR (4) | 10.3551 | 0.8329 | 3230811 | 2.9786 |
| Polynomial SVR (6) | 7.2376 | 0.6977 | 2298685 | 2.2241 |
| Gaussian SVR | 2.4222 | 0.5863 | 1124917 | 1.5115 |

Table 108: Results for R3,R4 \rightarrow R2-1000




Figure 108: Completion time vs ncores for: R3,R4 \rightarrow R2 with datasize 1000GB

S

6 Fixed Datasize, all the features,

6.1 Query Q2

6.1.1 Query Q2 — Datasize 40GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 1.6276 | -0.9999 | 32359 | 0.1341 |
| Linear SVR | 0.3547 | 0.9266 | 6650 | 0.0274 |
| Polynomial SVR (2) | 0.8078 | 0.6069 | 15924 | 0.0649 |
| Polynomial SVR (3) | 0.5883 | 0.7715 | 10921 | 0.0446 |
| Polynomial SVR (4) | 1.0508 | 0.1893 | 19767 | 0.0789 |
| Polynomial SVR (6) | 1.1443 | 0.0177 | 21436 | 0.0850 |
| Gaussian SVR | 0.5819 | 0.8864 | 10799 | 0.0422 |

Table 109: Results for Q2-40GB




Figure 109: Completion time vs ncores for query Q2 with datasize 40

6.1.2 Query Q2 — Datasize 50GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.7974 | 0.4045 | 26855 | 0.0982 |
| Linear SVR | 0.3127 | 0.9112 | 10937 | 0.0419 |
| Polynomial SVR (2) | 0.7898 | 0.4304 | 28540 | 0.1044 |
| Polynomial SVR (3) | 0.7750 | 0.4572 | 30331 | 0.1124 |
| Polynomial SVR (4) | 1.2471 | 0.2083 | 37425 | 0.1433 |
| Polynomial SVR (6) | 0.9696 | 0.2704 | 37310 | 0.1350 |
| Gaussian SVR | 0.5052 | 0.8391 | 17201 | 0.0663 |

Table 110: Results for Q2-50GB




Figure 110: Completion time vs ncores for query Q2 with datasize 50

6.2 Query Q3

6.2.1 Query Q3 — Datasize 50GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.1821 | 0.9385 | 5147 | 0.0238 |
| Linear SVR | 0.3589 | 0.8794 | 10758 | 0.0502 |
| Polynomial SVR (2) | 0.7194 | 0.0867 | 23497 | 0.1173 |
| Polynomial SVR (3) | 0.4076 | 0.7912 | 11214 | 0.0514 |
| Polynomial SVR (4) | 0.5708 | 0.5249 | 14385 | 0.0683 |
| Polynomial SVR (6) | 0.6185 | 0.4444 | 17186 | 0.0817 |
| Gaussian SVR | 0.1952 | 0.9502 | 5697 | 0.0269 |

Table 111: Results for Q3-50GB




Figure 111: Completion time vs ncores for query Q3 with datasize 50

6.3 Query Q4

6.3.1 Query Q4 — Datasize 50GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0990 | 0.9881 | 4688 | 0.0154 |
| Linear SVR | 0.1198 | 0.9849 | 6748 | 0.0217 |
| Polynomial SVR (2) | 0.7944 | 0.3575 | 53287 | 0.1731 |
| Polynomial SVR (3) | 0.4045 | 0.8300 | 25556 | 0.0869 |
| Polynomial SVR (4) | 0.7117 | 0.5033 | 42806 | 0.1298 |
| Polynomial SVR (6) | 0.7431 | 0.4922 | 45187 | 0.1372 |
| Gaussian SVR | 0.1867 | 0.9590 | 11207 | 0.0387 |

Table 112: Results for Q4-50GB




Figure 112: Completion time vs ncores for query Q4 with datasize 50

7 Fixed Datasize, only nCores,

7.1 Query Q2

7.1.1 Query Q2 — Datasize 40GB

| Model | RMSE | R^2 | Mean absolute error | Mean relative error |
|--------------------|--------|--------|---------------------|---------------------|
| Linear regression | 0.4732 | 0.7213 | 11516 | 0.0462 |
| Linear SVR | 0.4773 | 0.7465 | 10981 | 0.0436 |
| Polynomial SVR (2) | 0.9249 | 0.1892 | 21428 | 0.0832 |
| Polynomial SVR (3) | 0.5122 | 0.6857 | 11435 | 0.0460 |
| Polynomial SVR (4) | 0.9083 | 0.2025 | 21353 | 0.0831 |
| Polynomial SVR (6) | 0.9293 | 0.2194 | 21908 | 0.0848 |
| Gaussian SVR | 0.4556 | 0.7493 | 9715 | 0.0382 |

Table 113: Results for Q2-40GB




Figure 113: Completion time vs ncores for query Q2 with datasize 40

7.1.2 Query Q2 — Datasize 50GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.4165 | 0.8333 | 13809 | 0.0505 |
| Linear SVR | 0.4199 | 0.8411 | 14497 | 0.0514 |
| Polynomial SVR (2) | 1.0202 | 0.0000 | 38508 | 0.1365 |
| Polynomial SVR (3) | 0.5115 | 0.7521 | 16040 | 0.0601 |
| Polynomial SVR (4) | 1.0202 | 0.0000 | 38508 | 0.1365 |
| Polynomial SVR (6) | 1.0202 | 0.0000 | 38508 | 0.1365 |
| Gaussian SVR | 0.2887 | 0.9208 | 10435 | 0.0374 |

Table 114: Results for Q2-50GB




Figure 114: Completion time vs ncores for query Q2 with datasize 50

7.2 Query Q3

7.2.1 Query Q3 — Datasize 50GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.3813 | 0.7992 | 11506 | 0.0597 |
| Linear SVR | 0.3626 | 0.9342 | 10534 | 0.0556 |
| Polynomial SVR (2) | 0.8629 | -0.0000 | 27382 | 0.1376 |
| Polynomial SVR (3) | 0.4068 | 0.8520 | 11280 | 0.0585 |
| Polynomial SVR (4) | 0.8629 | -0.0000 | 27382 | 0.1376 |
| Polynomial SVR (6) | 0.8629 | -0.0000 | 27382 | 0.1376 |
| Gaussian SVR | 0.2917 | 0.9418 | 8764 | 0.0471 |

Table 115: Results for Q3-50GB




Figure 115: Completion time vs ncores for query Q3 with datasize 50

7.3 Query Q4

7.3.1 Query Q4 — Datasize 50GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.3026 | 0.9203 | 14054 | 0.0471 |
| Linear SVR | 0.3036 | 0.9208 | 14764 | 0.0491 |
| Polynomial SVR (2) | 1.0691 | 0.0077 | 71807 | 0.2703 |
| Polynomial SVR (3) | 0.4319 | 0.8457 | 25075 | 0.0836 |
| Polynomial SVR (4) | 1.0698 | 0.0077 | 71807 | 0.2722 |
| Polynomial SVR (6) | 1.0698 | 0.0077 | 71807 | 0.2722 |
| Gaussian SVR | 0.2880 | 0.9306 | 14033 | 0.0477 |

Table 116: Results for Q4-50GB




Figure 116: Completion time vs ncores for query Q4 with datasize 50

8 Fixed Datasize, all features with 1/nCores,

8.1 Query Q2

8.1.1 Query Q2 — Datasize 40GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 1.1724 | -0.0377 | 24035 | 0.1012 |
| Linear SVR | 0.3713 | 0.9310 | 7170 | 0.0293 |
| Polynomial SVR (2) | 0.7597 | 0.7219 | 15618 | 0.0647 |
| Polynomial SVR (3) | 0.5592 | 0.8204 | 10138 | 0.0414 |
| Polynomial SVR (4) | 0.9967 | 0.4517 | 18992 | 0.0756 |
| Polynomial SVR (6) | 1.1150 | 0.0733 | 20847 | 0.0824 |
| Gaussian SVR | 0.5455 | 0.9096 | 10367 | 0.0410 |

Table 117: Results for Q2-40GB




Figure 117: Completion time vs ncores for query Q2 with datasize 40

8.1.2 Query Q2 — Datasize 50GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.9401 | 0.1721 | 35276 | 0.1336 |
| Linear SVR | 0.4863 | 0.7826 | 15991 | 0.0631 |
| Polynomial SVR (2) | 0.7756 | 0.4662 | 28112 | 0.1025 |
| Polynomial SVR (3) | 0.7518 | 0.4994 | 26075 | 0.0987 |
| Polynomial SVR (4) | 1.1967 | 0.2245 | 36195 | 0.1382 |
| Polynomial SVR (6) | 1.1080 | 0.1845 | 37971 | 0.1412 |
| Gaussian SVR | 0.4857 | 0.8544 | 16546 | 0.0638 |

Table 118: Results for Q2-50GB




Figure 118: Completion time vs ncores for query Q2 with datasize 50

8.2 Query Q3

8.2.1 Query Q3 — Datasize 50GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.1821 | 0.9385 | 5147 | 0.0238 |
| Linear SVR | 0.3698 | 0.7939 | 9957 | 0.0472 |
| Polynomial SVR (2) | 0.6657 | 0.2370 | 22068 | 0.1086 |
| Polynomial SVR (3) | 0.4115 | 0.7344 | 10516 | 0.0487 |
| Polynomial SVR (4) | 0.6275 | 0.4275 | 16290 | 0.0778 |
| Polynomial SVR (6) | 0.6407 | 0.4513 | 16807 | 0.0790 |
| Gaussian SVR | 0.2121 | 0.9383 | 6227 | 0.0295 |

Table 119: Results for Q3-50GB




Figure 119: Completion time vs ncores for query Q3 with datasize 50

8.3 Query Q4

8.3.1 Query Q4 — Datasize 50GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.0990 | 0.9881 | 4688 | 0.0154 |
| Linear SVR | 0.1853 | 0.9602 | 11231 | 0.0364 |
| Polynomial SVR (2) | 0.6609 | 0.5640 | 41416 | 0.1327 |
| Polynomial SVR (3) | 0.3972 | 0.8135 | 23263 | 0.0776 |
| Polynomial SVR (4) | 0.4891 | 0.7383 | 28158 | 0.0915 |
| Polynomial SVR (6) | 0.6421 | 0.6013 | 39270 | 0.1209 |
| Gaussian SVR | 0.2128 | 0.9463 | 12910 | 0.0451 |

Table 120: Results for Q4-50GB




Figure 120: Completion time vs ncores for query Q4 with datasize 50

9 Fixed Datasize, only 1/nCores,

9.1 Query Q2

9.1.1 Query Q2 — Datasize 40GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.4755 | 0.7186 | 9942 | 0.0394 |
| Linear SVR | 0.4356 | 0.7670 | 9580 | 0.0378 |
| Polynomial SVR (2) | 0.7545 | 0.3885 | 17726 | 0.0705 |
| Polynomial SVR (3) | 0.5410 | 0.6751 | 12134 | 0.0488 |
| Polynomial SVR (4) | 0.6630 | 0.5285 | 15720 | 0.0630 |
| Polynomial SVR (6) | 0.6016 | 0.5722 | 14553 | 0.0587 |
| Gaussian SVR | 0.4536 | 0.7500 | 9672 | 0.0381 |

Table 121: Results for Q2-40GB




Figure 121: Completion time vs ncores for query Q2 with datasize 40

9.1.2 Query Q2 — Datasize 50GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.3294 | 0.8958 | 11482 | 0.0409 |
| Linear SVR | 0.3301 | 0.8973 | 11520 | 0.0410 |
| Polynomial SVR (2) | 0.8382 | 0.4016 | 32740 | 0.1153 |
| Polynomial SVR (3) | 0.4515 | 0.8074 | 15713 | 0.0568 |
| Polynomial SVR (4) | 0.6385 | 0.6123 | 21200 | 0.0775 |
| Polynomial SVR (6) | 0.5805 | 0.6807 | 18767 | 0.0683 |
| Gaussian SVR | 0.2887 | 0.9208 | 10434 | 0.0374 |

Table 122: Results for Q2-50GB




Figure 122: Completion time vs ncores for query Q2 with datasize 50

9.2 Query Q3

9.2.1 Query Q3 — Datasize 50GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.4364 | 0.7371 | 14301 | 0.0707 |
| Linear SVR | 0.3581 | 0.8593 | 10636 | 0.0527 |
| Polynomial SVR (2) | 0.7692 | 0.1922 | 22882 | 0.1221 |
| Polynomial SVR (3) | 0.4719 | 0.7356 | 12895 | 0.0635 |
| Polynomial SVR (4) | 0.6625 | 0.4085 | 18813 | 0.0980 |
| Polynomial SVR (6) | 0.6111 | 0.5019 | 16636 | 0.0864 |
| Gaussian SVR | 0.2815 | 0.9430 | 8413 | 0.0456 |

Table 123: Results for Q3-50GB




Figure 123: Completion time vs ncores for query Q3 with datasize 50

9.3 Query Q4

9.3.1 Query Q4 — Datasize 50GB

| Model | RMSE | R ² | Mean absolute error | Mean relative error |
|--------------------|--------|----------------|---------------------|---------------------|
| Linear regression | 0.4439 | 0.8284 | 26027 | 0.0969 |
| Linear SVR | 0.4490 | 0.8300 | 27044 | 0.1029 |
| Polynomial SVR (2) | 1.0085 | 0.1628 | 66229 | 0.2716 |
| Polynomial SVR (3) | 0.6029 | 0.7055 | 37412 | 0.1473 |
| Polynomial SVR (4) | 0.8773 | 0.3698 | 55542 | 0.2300 |
| Polynomial SVR (6) | 0.7938 | 0.4631 | 47441 | 0.1942 |
| Gaussian SVR | 0.2888 | 0.9308 | 14254 | 0.0484 |

Table 124: Results for Q4-50GB




Figure 124: Completion time vs ncores for query Q4 with datasize 50